

ZF-Microglia-AI — Software Guide

A developer/technical reference: every module, every function, what algorithm it implements, and how the pieces connect.

Companion to [GUIDE.md](#) (user-facing, how to operate the plugin) and [STATISTICS GUIDE.md](#) (the statistics formulas in detail — this guide summarizes them, that one derives them). Written for someone reading the source for the first time, or citing the pipeline’s algorithms in a thesis/paper.

Pseudocode blocks use a plain, language-agnostic notation (not literal Python) — `for/if/while`, indentation for scope, `=` for assignment where it helps distinguish from equality, and function calls written `name(args)`. Real function/parameter names are used throughout so a reader can go from pseudocode straight to `grep` in the real source.

Table of Contents

1. [Architecture overview](#)
2. [Common infrastructure](#) — file I/O, secrets, GPU detection, live-progress capture, generic coarse-to-fine sweeping, `_widget.py`’s own shared helpers
3. [Stage 1 — Skin removal](#) — `_inference.py`, `_background.py`
4. [Stage 2 — Label creation](#) — `_labeling.py` (Pixel Classifier), `_cellpose_seg.py` (Cellpose-SAM), `_sanding.py`
5. [Stage 3 — Label editing & correction](#) — `_labeling.py` (editing tools), `_grow_correct.py`, `_contrast_sweep.py`, `_auto_correction.py`
6. [Stage 4 — Statistics](#) — `_statistics.py`
7. [Stage 5 — AI Tools](#) — `_gt_annotation.py`, `_xxyz_patches.py`, `_crop_truncation.py`, `_branch_calibration.py`, `_training_jobs.py`
8. [Stage 6 — Sweeps & Utilities](#) — `_sieve.py`, `_brain_sweep.py`, `_pixel_sweep.py`, `_krendl_sweep.py`, `_epoch_sweep.py`, `_gt_score.py`, `_gt_package.py`, `_gt_toolkit.py`
9. [_widget.py — application shell](#)
10. [Module reference \(quick index\)](#)

1. Architecture overview

What the plugin does, as a pipeline

```
raw confocal stack (.tif/.ims)
    | Stage 1 – Skin Removal (_inference.py, _background.py)
    ▼
brain_only volume (_ExtRm / _NoBG / _RndFill)
    | Stage 2 – Label Creation (_labeling.py or
_cellpose_seg.py)
    ▼
Labels layer (one integer ID per cell)
    | Stage 3 – Label Editing & Correction (_labeling.py,
_grow_correct.py,
    | _contrast_sweep.py, _auto_correction.py)
    ▼
corrected Labels layer
    | Stage 4 – Statistics (_statistics.py)
    ▼
per-cell CSV (shape, branching, spatial, intensity, optional LLM
description)
```

Two more subsystems sit alongside this straight-line pipeline, not inside it:

- **Stage 5 — AI Tools:** builds and trains the two models (Stages 1 and 2 depend on (MONAI for skin removal, Cellpose-SAM for label creation), plus the ground-truth annotation tool that produces the training data in the first place.
- **Stage 6 — Sweeps & Utilities:** GT-verification tools that check a pipeline parameter against real ground truth and (when confirmed) recalibrate it — this is how the “recommended” defaults quoted throughout GUIDE.md were actually found, not guessed.

Where state lives

There is no persistent application database. Three places hold state:

1. **napari’s own layer list** (`self._viewer.layers`) — the volumes and label arrays themselves, as `napari.layers.Image/napari.layers.Labels` objects. Every pipeline stage reads its input

from a layer and writes its output as a new (or in-place-mutated) layer. This is the plugin's real "shared memory" between stages.

2. `~/ .config/napari-zf-microglia-ai/config.json`

(`self._state["config"]`, read/written via `_load_config()`/`_save_cfg()` in `_widget.py`) — every slider/checkbox default, every "recommended" GT-swept value (see `_update_gt_history()`), file-browse paths, and email/API settings (except the secrets themselves — see `_secrets.py` in §2).

3. **The OS credential store or a local encrypted file** (`_secrets.py`) — the two values deliberately kept out of `config.json`: the SMTP app password and the LLM API key.

The threading pattern used by almost every "Run"/"Sweep" button

Every long operation in this plugin — MONAI inference, Cellpose-SAM's `do_3D`, a GT sweep, a training launch — follows the same shape, to keep napari's own Qt event loop responsive:

```
def _on_run_button_click():
    validate inputs; if invalid → show error, return
    disable the button; set status text to "Running..."
    result ← {} # plain dict, mutated
    by the worker
    def _worker():
        try:
            result["value"] ← do_the_actual_work(...) # runs on
a background thread
        except Exception as exc:
            result["error"] ← str(exc)
    thread ← Thread(target=_worker); thread.start()

    timer ← QTimer() # polls on the GUI
thread
    def _poll():
        if thread.is_alive():
            if "_progress" in result: update status text
            return # check again next
tick
        timer.stop()
        if "error" in result: show it; re-enable button; return
        apply result["value"] to the relevant napari layer(s) #
GUI-thread-only work
        re-enable button
```

```
timer.timeout.connect(_poll); timer.start(200)    # re-check
every 200 ms
```

`result` is the hand-off point between the background thread (which must never touch Qt widgets or `viewer.layers` directly — napari/Qt are not thread-safe) and the GUI thread (which owns every widget and layer mutation). A `progress_cb` callback threaded through the actual algorithm lets the worker report intermediate status without touching the GUI itself — it just writes into `result["_progress"]`, which the next poll tick picks up.

Two further patterns build on this:

- **Detached processes** (`_training_jobs.py`, Stage 5): a training run is a genuinely separate OS process (`conda run ...`), not a background thread inside napari — it survives napari closing entirely. The same poll-a-shared-state loop just polls a PID + log file on disk instead of an in-process Thread.
- **The hang watchdog** (`_hang_watch_start/_hang_watch_stop` in `_widget.py`, §9): wraps the thread-start of the slowest operations (Protect Skin as Label, Auto-correct) with `faulthandler.dump_traceback_later()`, so a run that's still going after 30 minutes prints every thread's real stack to the terminal — turning “is it stuck?” into “here's exactly which line it's on.”

Layer-selection philosophy

Early in this project every editing tool read its input from whatever layer happened to be “active” in napari's own layer list (`_active_labels_layer()`: active selection, else topmost Labels layer, else None) — implicit, and silently wrong whenever the wrong layer happened to be selected. Every tool in Stage 3 (Edit MG Labels) now instead reads from one **explicit, shared combo-box selector** (`_edit_labels_layer()`, `self._edit_signal_combo/_edit_labels_combo/_edit_mask_combo`) that the user sets once, not per tool — see §9.

2. Common infrastructure

Modules and helpers used across more than one pipeline stage.

_io.py — file loading with metadata-aware voxel scale

Every downstream algorithm that measures a physical distance (μm) needs the true voxel size (Z spacing vs. XY pixel size — these stacks are anisotropic, typically $\sim 5.75:1$). This module's only job is finding that scale and loading the raw pixel data consistently for .tif/.tiff/.ims.

_calc_anisotropy(voxel_z_um, voxel_x_um, voxel_y_um)

```
xy ← (voxel_x_um + voxel_y_um) / 2
return voxel_z_um / xy    (or 1.0 if xy ≤ 0)
```

extract_tif_metadata(tif_path) — reads ImageJ-embedded metadata from a TIFF's own tags (spacing/finterval for Z, XResolution/YResolution for XY, which ImageJ stores as *pixels per μm* , hence the reciprocal). Returns None on any failure so the caller can fall through to the next metadata source.

parse_metadata(metadata_path) — parses a Leica-format *_metadata.txt sidecar file with regexes for camera pixel width/height, Z step size, and every TotalConsolidatedOpticalMagnification value present (there are usually two: the objective's, $\geq 10\times$, and a digital zoom factor, $< 10\times$):

```
cam_x, cam_y ← regex-extracted "Pixel Width/Height ( $\mu\text{m}$ )"
voxel_z      ← regex-extracted "StepSize"
mags         ← every "TotalConsolidatedOpticalMagnification" value
              found
obj_mag      ← first mag ≥ 10           (the objective)
zoom_mag     ← first mag < 10          (digital zoom)
total_mag    ← obj_mag × zoom_mag
voxel_x, voxel_y ← cam_x / total_mag, cam_y / total_mag
```

Returns None if any of the four regex groups is missing.

find_best_metadata_match(file_path) — a stack's own metadata file isn't always named exactly <stem>_metadata.txt; this looks for an exact match first, then the closest name (difflib.SequenceMatcher string-similarity ratio, threshold 0.3) among every *_metadata.txt in the same folder, then repeats one directory up if nothing was found (metadata sometimes lives one level above the stack).

load_file(path) — the public entry point:

```
if suffix is .tif/.tiff:
    raw ← tifffile.imread(path)
```

```

        metadata ← find_best_metadata_match(path) →
parse_metadata(...)
        or extract_tif_metadata(path)
        or the hardcoded default (1,1,1) μm/vox
    if raw is 3D: channels ← [(raw, stem, metadata)]
    if raw is 4D: channels ← [(raw[c], f"{stem}_ch{c}", metadata)
for each channel c]
elif suffix is .ims:
    f ← ImsReader(path)
    metadata ← find_best_metadata_match(...) → parse_metadata(...)
        or f's own embedded resolution
    channels ← [(f.get_Volume_At_Specific_Resolution(...), name,
metadata)
        for each channel]
return channels    # list of (volume, layer_name, metadata_dict)

```

Every channel is returned separately so the caller (`_reader.py`, or `_widget.py`'s own Open-file handler) can add each as its own napari Image layer, letting the user pick which channel is the actual microglia signal.

`_reader.py` — napari File-menu integration

Registers this plugin as a napari reader (nape2/napari.yaml “readers” contribution) so `.tif/.ims` files opened via napari’s own File → Open work correctly, not just via this plugin’s own Open button.

```

get_reader(path):
    if path's suffix in { .tif, .tiff, .ims }: return _read_file
    else: return None                                # tells napari "not
mine"

_read_file(path):
    channels ← _io.load_file(path)
    for each (volume, name, metadata) in channels, with index i:
        colormap ← _CH_COLORMAPS[i mod 4]           # gray, green,
magenta, cyan, repeating
        yield (volume, {name, colormap, scale: metadata["scale"]},
"image")

```

`_secrets.py` — layered credential storage

Two values need to persist between sessions without ever sitting in the plugin's own plaintext `config.json`: the SMTP app password (email notifications) and the LLM API key (Statistics' description backend).

Two-tier storage:

```
get_secret(key):
    try: value ← keyring.get_password(SERVICE_NAME, key)    # Tier
1: OS credential store
        if value: return value
    except: pass
    return _fallback_get(key)                                # Tier
2: local Fernet-encrypted file

set_secret(key, value):
    try:
        keyring.set_password/delete_password(...)          # Tier
1
        _fallback_set(key, "")    # clear any stale Tier-2 copy
now that Tier 1 works
        return None    # success, nothing to report
    except Exception as tier1_exc:
        try:
            _fallback_set(key, value)                        #
Tier 2
            return None (but print a warning naming the weaker
guarantee)
        except Exception as tier2_exc:
            return an error string (both tiers failed -- caller
shows it, doesn't crash)
```

`_get_fernet()` generates a machine-local AES key on first use (`Fernet.generate_key()`), stores it `chmod 600` next to the encrypted blob — a real but deliberately weaker guarantee than Tier 1 (the key lives beside what it protects), kept because Tier 1 (keyring's Secret Service backend) reliably fails on SSH-only Linux sessions with no PAM-driven keyring unlock, confirmed directly on this project's own workstation.

`migrate_plaintext_secrets(cfg)` — a one-time upgrade path: if an old `config.json` still has `api_key/notify_smtp_password` in plaintext (from before this module existed), move them into `set_secret()` and strip them from the dict that gets written back.

`_gpu_check.py` — GPU/VRAM classification

Checked once at import time (hardware doesn't change mid-session) and cached into module-level constants (GPU_HAS_CUDA, GPU_VRAM_GB, GPU_MEETS_RECOMMENDED, GPU_NAME, GPU_MSG):

```
check_gpu():
    try:
        if not torch.cuda.is_available(): return {has_cuda:
False, ...}
        props ← torch.cuda.get_device_properties(0)
        vram_gb ← props.total_memory / 10243
        meets ← vram_gb ≥ 8
        return {has_cuda: True, vram_gb, meets_recommended: meets,
name: props.name}
    except: return {has_cuda: False, ...}    # never raises --
worst case, a cautious banner
```

Informational only — Tab 5 (AI Tools) is always shown regardless of the result; this only decides which color/text its disclaimer banner uses.

`_live_progress.py` — surfacing libraries' own hidden progress output

MONAI's `sliding_window_inference(progress=True)` draws a `tqdm` bar to `raw sys.stderr`; Cellpose routes everything through Python's logging module but never calls its own `logger_setup()` unless its CLI is used, so by default its log records go nowhere — not even a terminal. `capture_live_output(push)` is a context manager that makes both visible in the GUI for the duration of one call:

```
capture_live_output(push):
    if another capture is already active: yield (no-op
passthrough); return
    save real sys.stdout/sys.stderr and the root logger's level
    sys.stdout ← sys.stderr ← a _LineSink(push)    # buffers
partial writes,
                                                    # pushes a
complete line on \n OR \r
                                                    # (tqdm
redraws with \r; GUI can't
                                                    # overwrite
in place, so \r → new line)
```

```

        attach a logging.Handler to the root logger that calls
push(formatted record)
        set root logger level to INFO if it was less verbose
        yield                                     # caller's real
work runs here
        finally: restore stdout/stderr, remove the handler, restore
the log level

push must be thread-safe (the caller always passes the same
result["_progress"] = msg pattern every other worker uses).

```

`_sieve.py` — generic coarse-to-fine sweep narrowing

Shared by every GT-sweep tool that has one continuous “primary axis” (MONAI Threshold, BG Threshold, Smooth σ XY, Cellprob) — born from a real manual workflow (sweep a wide range coarsely, narrow around the winner, narrow again) generalized into one reusable engine that knows nothing about what it’s sweeping.

```

sieve_grid(lo, hi, center, half_width, step):
    stage_lo ← max(lo, center - half_width)
    stage_hi ← min(hi, center + half_width)
    if stage_hi ≤ stage_lo: return [clip(center, lo, hi)]      #
degenerate range → 1 point
    return every value from stage_lo to stage_hi, step apart

run_sieve(run_stage_fn, lo, hi, coarse_step, refine_steps,
progress_cb, cancel_check):
    stage_specs ← [(None, None)] + refine_steps      # stage 1 =
full range, no narrowing
    values ← grid from lo to hi, step = coarse_step
    for each (half_width, step) in stage_specs, with index i:
        if cancel_check(): mark cancelled; break
        if i > 0: values ← sieve_grid(lo, hi, best_value,
half_width, step)
        best_value, stage_result ← run_stage_fn(values)      # caller
runs ONE 1D grid,
                                                                    #
everything else held fixed
        record this stage (values, best_value, stage_result)
        if stage_result says cancelled: mark cancelled; break
    return {stages: [...], final_best_value, final_result,
cancelled}

```

`run_stage_fn` is the only thing that actually knows about MONAI/Cellpose/etc. — it’s a closure the caller builds, e.g. “run the BG-Threshold sweep at these candidate values, holding Erosion fixed at its current slider value, return the winning threshold and the sweep’s own report dict.”

`_widget.py`’s own shared helpers

A handful of small functions used throughout `_build_ui()` (the ~10,000-line method that constructs every tab — see §9), not tied to any one pipeline stage:

- **`_add_reliable_spinbox(row_layout, slider, min, max, step, decimals=None)`** — every numeric slider in this plugin is paired with a real `QSpinBox/QDoubleSpinBox` synced bidirectionally (`slider.valueChanged` → `spin.setValue`, `spin.valueChanged` → `slider.setValue`), replacing `superqt`’s own built-in numeric label whose width reset unpredictably on first show. Returns the `spinbox`.
- **`_set_layout_widgets_visible(layout, visible)`** — recursively walks a `QVBoxLayout/QHBoxLayout` tree (rows mix `addWidget` and nested `addLayout` calls) and calls `setVisible()` on every real widget found, since only widgets (not layouts) support it directly.
- **`_make_collapsible(groupbox, start_expanded=True)`** — turns any already-built `QGroupBox` into a collapsible section: makes its own title bar checkable, and toggling it calls `_set_layout_widgets_visible` on the group’s contents. Every group box in every tab goes through this.
- **`_wrap_scroll(widget)`** — wraps a fully-built tab page in a vertical-only `QScrollArea` (`setWidgetResizable(True)`, horizontal scrollbar forced off) so a tab taller than the napari window stays fully reachable.
- **`_add_recommended_label(layout, initial, unit, noun)`** — the read-only “Recommended: X” line under a GT-sweepable slider, kept separate from the slider’s own editable value so a GT-confirmed sweep’s finding is never silently overwritten by the user nudging the slider to test something else.
- **`_add_gt_checkbox(layout, hint, visible=True)`** — the “This is verified ground truth” one-shot opt-in every sweep tool has: off by default, un-checks itself after every run, and is the only thing that lets a sweep move a shared recommended value (Min volume, a slider’s own “Recommended” line, etc.) rather than just reporting a result.

- **_make_notify_checkbox()** — builds one “Email me when done” QCheckBox with a shared tooltip; every long-running, non-detached tool (Run Skin-Remover, Cellpose-SAM Segmentation, the Cellprob/Best-Epoch sweeps) gets its own instance of this, all reading the same shared SMTP credentials configured once in Tab 6.
-

3. Stage 1 — Skin removal

UI: Tab 1 — Skin Remover. **Modules:** `_inference.py` (MONAI brain segmentation), `_background.py` (three background-removal modes).

Goal: given a raw confocal stack, produce `brain_only` — everything but the brain zeroed — for Stage 2 to label.

`_inference.py` — MONAI 3D U-Net sliding-window inference

`_normalize(volume)` — percentile normalization, the same transform the model was trained under:

```
p01, p99 ← 1st and 99th percentile of volume
return clip((volume - p01) / max(p99 - p01, 1), 0, 1)
```

`predict_probability(volume, model_path, device)` — the expensive half of inference, split out on its own so a GT-sweep tool can run it once and cheaply re-threshold many candidate values afterward (see §8):

```
model ← UNet(3D, 1→1 channel, channels=(32,64,128,256,512), 2 res-
units/level)
    load weights from model_path checkpoint; eval mode
volume_t ← normalize(volume) as a (1,1,Z,Y,X) tensor on device
pred_logits ← sliding_window_inference(volume_t, roi=(64,192,192),
    batch=4, overlap=0.25, mode="gaussian",
predictor=model)
    # slides a 64×192×192 window across the whole
volume, blending
    # overlapping predictions with a Gaussian-weighted
average
    # (softer seams than a hard-edge tile stitch)
return sigmoid(pred_logits)    # per-voxel probability, 0..1
```

`postprocess_probability(pred_prob, threshold)` — the cheap half, safe to call repeatedly per candidate threshold:

```

raw_mask <- pred_prob > threshold
labeled, n_components <- connected_components(raw_mask)
clean <- the single largest component          # drops isolated
blobs / mislabeled tissue
clean <- fill_holes(clean)                    # solid,
contiguous brain
return clean as uint8    # "brain_mask", always un-eroded

```

run_inference(volume, model_path, threshold, device, erosion_voxels) — glues the two together and applies the optional erosion:

```

pred_prob <- predict_probability(...)
brain_mask <- postprocess_probability(pred_prob, threshold)    #
always un-eroded – this is

# what gets saved as brain_mask.tif
if erosion_voxels > 0:
    eroded_mask <- binary_erosion(brain_mask,
iterations=erosion_voxels)
else:
    eroded_mask <- brain_mask
brain_only <- volume × eroded_mask          # zero everything the
(possibly eroded) mask excludes
return brain_mask, brain_only, eroded_mask

```

eroded_mask is the one downstream background-removal steps must use as “the brain boundary” — using brain_mask there instead would silently discard the Erosion slider.

`_background.py` — background estimation and removal, 3 modes

All three modes share one background estimate and one threshold rule:

`_estimate_background(volume, brain_mask):`

```

brain_pixels <- volume[brain_mask]          # only INSIDE
the (possibly eroded) brain
hist <- 1000-bin histogram of brain_pixels
bg_mode <- the histogram bin center with the highest count    # the
population mode, robust

#
to a skewed intensity tail

```

```

bg_values ← every brain_pixel ≤ bg_mode
return bg_values, bg_mode, bg_mode, bg_mode          # (median/
min/max all collapse to bg_mode                      # in the
current implementation)

```

_threshold(volume, brain_mask, tolerance_pct):

```

bg_max ← _estimate_background(...)'s bg_mode
tol ← (max(volume) - min(volume)) × tolerance_pct / 100
threshold ← bg_max + tol
bg_mask ← volume ≤ threshold          # single-sided:
everything at/below threshold = background
return bg_values, bg_max, threshold, bg_mask

```

Mode 1 — remove_outside_brain (_ExtRm suffix): zero background pixels *outside* the brain only; everything inside the brain mask is kept untouched regardless of intensity.

```

bg_mask ← _threshold(...)
to_zero ← (~brain_mask) AND bg_mask
result ← volume.copy(); result[to_zero] ← 0

```

(Note: as actually wired in _widget.py's _on_run, Mode 1's saved brain_only is volume × eroded_mask directly — the entire exterior zeroed, not remove_outside_brain's more selective “background-only” result — see §9.)

Mode 2 — remove_global (_NoBG suffix): zero every background-threshold pixel in the *whole* stack, inside the brain included — this is the input the Pixel Classifier route needs (background must be zero everywhere for its own threshold/union-find pipeline to find clean isolated blobs).

```

bg_mask ← _threshold(...)
if signal_erosion_voxels > 0:
    signal_mask ← erode_signal_2d(~bg_mask,
signal_erosion_voxels)  # shrink what SURVIVES
    bg_mask ←
~signal_mask          #
thresholding, not the mask
result ← volume.copy(); result[bg_mask] ← 0

```

erode_signal_2d(mask, iterations) erodes one Z-slice at a time (binary_erosion per 2D slice), not a single 3D erosion — a plain 3D erosion also eats iterations slices off the top *and* bottom of a blob's own Z-extent, which for a thin blob can delete it outright; per-slice erosion only ever shrinks each slice's in-plane contour.

Mode 3 — fill_outside_brain_random (_RndFill suffix, presentation-only, not fed into any labeling route):

```
bg_mode ← _estimate_background(...)
tol ← 0.001 × (max(volume) - min(volume))          # ±0.10% fill
window around the mode
result ← volume.copy()
result[outside brain] ← uniform_random(bg_mode - tol, bg_mode +
tol), one value per outside voxel
```

_widget.py — Tab 1 handlers

_on_open() — background-thread wrapper around
_io.load_file(path); on completion, adds every returned channel as
its own Image layer (_add_channels) and reports channel count + shape.

_on_load_labels() — loads a saved .tif directly as a Labels layer:
reads with tiffio.imread(...).astype(int32), scales it to match
whatever stack is already open (self._state["metadata"]["scale"],
warning if none is open yet), and names it <active Image
layer>_labels — the same convention Stage 2's own label-creation
tools use, so a reloaded file fits the rest of the plugin's naming scheme.

_on_run() — the main button:

```
validate: model file exists, an Image layer is selected and 3D
read Threshold/Erosion/Background-mode/BG-Threshold/Signal-Erosion
from their sliders
device ← cuda, else mps, else cpu
background thread, wrapped in capture_live_output (surfaces
MONAI's tqdm bar +
    _background.py's own print() lines into the GUI's log view,
not just a terminal):
    brain_mask, brain_only, eroded_mask ← run_inference(volume,
model_path, threshold,
                                                    device,
erosion_voxels)
    if bg_mode == 1 (_ExtRm): brain_only ← volume × eroded_mask
    if bg_mode == 2 (_NoBG): brain_only ← remove_global(volume,
```

```

eroded_mask, ...) × eroded_mask
    if bg_mode == 3 (_RndFill): brain_only ←
fill_outside_brain_random(volume, eroded_mask)
    # bg_mode == 0 (Off): brain_only stays exactly
run_inference()'s own (volume × eroded_mask)
on completion: add brain_mask/brain_only as new Labels/Image
layers (suffix per _BG_SUFFIX),
    save the ticked output files, send email notification if
configured

```

4. Stage 2 — Label creation

UI: Tab 2 — Create MG Labels. Two independent routes, auto-shown/hidden by the active layer’s background-mode suffix (`_NoBG` → Pixel Classifier, `_ExtRm` → Cellpose-SAM): **Pixel Classifier** (`_labeling.py`’s `create_labels`, a plain 3D connected-components engine) and **Cellpose-SAM** (`_cellpose_seg.py`’s `run_full_pipeline`, a trained-network `do_3D` segmentation + several correction stages). Both finish with an optional **Sanding** pass (`_sanding.py`). **Goal:** turn `brain_only` into a Labels layer, one integer ID per cell.

`_labeling.py` — Pixel Classifier: true 3D connected components

Despite older project history calling this a “union-find” engine, the current implementation is a plain 3D connected-components labeling, not a per-slice-then-stitch union-find — `create_labels()` dispatches to whichever of 3 backends is fastest (detected once at import time: **CUDA** via `CuPy/cupyx` > **Apple MPS** → falls back to threaded CPU, since MPS itself has no `ndimage` ops > **CPU**, `scipy.ndimage` + a `ThreadPoolExecutor` at 50% of cores). All three backends implement the identical 6-step algorithm; only the array library and slice-loop parallelism strategy differ.

```

create_labels(volume, sigma_xy, sigma_z, min_volume,
min_hole_size, final_min_fraction):

```

```

1. binary ← volume > 0
2. blurred ← gaussian_filter(binary, sigma=(sigma_z, sigma_xy,
sigma_xy))
    smooth_mask ← blurred > 0.5           # re-thresholded smooth
binary mask —

```

```

# this is what actually
connects/

# disconnects nearby
blobs, not the

# raw thresholded signal
3. for each Z slice:
    fill any enclosed background hole whose area <
min_hole_size
    (min_hole_size <= 0 → fill every hole unconditionally, the
    original behaviour)
4. labeled, n_objects ← connected_components(smooth_mask,
structure=ones(3,3,3))
    # 26-connectivity: any of the 26
neighbors counts
5. threshold ← max(1, round(final_min_fraction × min_volume))
    # NOT min_volume itself – same golden-
ratio
    # safety-net relaxation
_cellpose_seg.py's
    # final_min_size_cleanup() uses, since
this route
    # has no later merge/reattach stage to
give a
    # gray-zone object a second chance
    remove every blob smaller than threshold
6. renumber surviving blobs 1..N by DESCENDING volume (label 1 =
largest)
return int32 labels, 0=background

```

The CUDA path (`_create_labels_cuda`) runs steps 1–2 and the volume/threshold math as vectorized CuPy array ops (`cp.bincount` for per-label voxel counts instead of a Python loop), loops step 3 per Z slice on the GPU, and frees the CuPy + PyTorch memory pools (`_free_gpu_cache()`) before returning — including on the exception path, so a mid-run CUDA OOM still falls back to CPU cleanly (`create_labels()`'s own try/except around the CUDA call). The threaded-CPU path (`_create_labels_threaded`, also used for Apple MPS) runs step 3's per-slice hole fill via a `ThreadPoolExecutor`, everything else as a single `scipy.ndimage/numpy` call over the whole volume.

min_hole_size mechanism (`_fill_holes_capped_gpu` / the CPU per-slice closure): unconditional `binary_fill_holes` fills every enclosed background region no matter its size — much like Cellpose's own default hole-filling (see §4's Cellpose-SAM section), a single-voxel prediction artifact and a genuine internal structural void look identical

to it. A positive `min_hole_size` switches to an area-limited fill (`skimage.morphology.remove_small_holes` on CPU; a manual `label() + bincount()`-based size filter on GPU, since `cupyx` has no equivalent), leaving any hole at or above the floor standing as real background.

`_cellpose_seg.py` — Cellpose-SAM: do_3D + 4-stage cleanup pipeline

Same math as the project's earlier standalone `krend1_do3d.py` CLI script, refactored into importable functions. No GT-based relabeling/scoring here (that stays a CLI/research workflow, see §8) — this produces a clean, sequentially-labeled instance mask ready for manual correction.

`predict_flows(volume, model_path, anisotropy, gpu)` — the one genuinely GPU-bound, expensive step: a single Cellpose-SAM network forward pass producing a per-voxel flow field (`dP`) and cell-probability map (`cellprob`), with neither `cellprob_threshold` nor `flow_threshold` applied yet — both are cheap, CPU/light-GPU post-processing steps applied afterward. Split out on its own (mirroring `_inference.py`'s `predict_probability/postprocess_probability` split for MONAI) so a `cellprob` sweep can call this once and cheaply re-threshold many candidate values without repeating the expensive network pass:

```
model ← CellposeModel(pretrained_model=model_path, gpu=gpu)
_, flows, _ ← model.eval(volume, do_3D=True,
    anisotropy=anisotropy, z_axis=0,
                                channel_axis=None, diameter=None,
    normalize=True,
                                augment=False, compute_masks=False) #
compute_masks=False
#

skips mask formation
dP, cellprob ← flows[1], flows[2]
return model, dP, cellprob, volume.shape
```

`save_flow_cache / load_flow_cache` — persist `predict_flows()`'s expensive output to disk (`np.savez_compressed`) plus a fingerprint (model checkpoint path, anisotropy, volume shape/dtype). Exists because mask formation (`masks_from_flows`'s own internal `follow_flows()` re-run) needs its own separate chunk of GPU memory and can CUDA-OOM *even after* a multi-hour network pass already succeeded — without a cache, that OOM silently destroys the entire network pass (Python's own except-block variable cleanup, confirmed by reading `_widget.py`'s own handler, keeps only `str(exc)`, nothing

else). `run_full_pipeline()`'s `flow_cache_path` saves the cache right before mask formation starts and deletes it automatically the moment mask formation succeeds — it only survives on disk if that specific step crashes, and a stale/mismatched cache (different fish/model/anisotropy) is detected via the fingerprint and ignored, never silently reused.

`masks_from_flows(model, dP, cellprob, shape, cellprob_threshold, flow_threshold, min_size, max_size_fraction, niter, min_hole_size, progress_cb)` — cheap: forms instance masks from an already-computed flow field. `do_3D=True` is baked in. `flow_threshold` is accepted only to match `do_3D`'s call signature — Cellpose's own `compute_masks()` only applies its flow-error QC filter when `do_3D=False` (confirmed directly in `cellpose/dynamics.py`; a call-count spy test found 0 calls under `do_3D=True`), so it's a documented no-op here. `niter=None` is resolved to Cellpose's own default of 200 (calling `_compute_masks()` directly, as this function does, skips `eval()`'s own resolution of that default, which would otherwise crash with a `TypeError` on `range(None)`).

```
if niter is None: niter ← 200
report progress (no further progress prints until this whole step
finishes –
    neither this function nor Cellpose's own
follow_flows()/
    compute_masks() log anything while it runs)
monkey-patch cellpose.utils.fill_holes_and_remove_small_masks →
    _make_capped_fill_holes(min_hole_size)    # see below
try:
    masks ← model._compute_masks(shape, dP, cellprob,
    flow_threshold, cellprob_threshold,
                                min_size, max_size_fraction,
niter, do_3D=True)
finally:
    restore the original fill_holes_and_remove_small_masks
return masks as int32
```

`_make_capped_fill_holes(min_hole_size)` — builds a drop-in, size-aware replacement for Cellpose's own hole-filling (which unconditionally calls `fill_voids.fill()` over each predicted mask's full 3D crop, no size threshold at all — the same category of bug `create_labels()`'s own `min_hole_size` was built to fix). `min_hole_size≤0` keeps Cellpose's exact original unconditional behavior. A positive value switches to a **per-Z-slice** `skimage.morphology.remove_small_holes` loop, not one 3D call over the whole crop — a real per-slice-small hole that

persists across many Z-slices is one 3D-connected void whose *total* volume can dwarf the threshold, so a plain 3D `remove_small_holes()` judges it “too big to be noise” and leaves a visible ring on every slice it spans; looping per-slice avoids that (mirrors `_labeling.py`’s own per-slice treatment).

`gmm_cleanup(masks)` — 3-component Gaussian Mixture Model on $\log_{10}(\text{volume})$ of every raw predicted object, separating a noise / gray-zone / real-cell population and dropping everything below the auto-detected gray→cell intersection point:

```
if fewer than 3 objects: return unchanged (not enough data to fit
3 components)
x ←  $\log_{10}$ (each object's voxel volume)
gmm ← GaussianMixture(n_components=3).fit(x)
sort the 3 fitted components by mean: small, mid, large
cutoff ← the point where the mid and large components' fitted
Gaussian
           curves cross (closed-form solution of two Gaussians'
log-density
           equality; falls back to the midpoint of the two means if
the
           quadratic has no real root in range)
remove every object smaller than cutoff
return masks, cutoff, n_removed
```

`krendl_safe_merge(masks, max_gap, min_contact, gt_min, scale_zyx)`
— merges only genuinely small fragments ($< \text{gt_min}$, the smallest true GT-labeled cell volume ever confirmed) into their nearest larger neighbor, when either close enough or touching with enough contact area:

```
repeat up to 200 passes:
    candidates ← every fragment below gt_min, smallest first
    for each candidate fragment f:
        find the nearest larger-volume object t (bbox-proximity-
prefiltered
        in voxel space, using max_gap converted to a voxel margin
via the
        FINEST axis – safe over-inclusion, never wrongly rejects a
close pair)
        gap ← min distance from f's mask to t's mask, via a
physically-scaled
        EDT (distance_transform_edt(~f_mask,
sampling=scale_zyx)) –
```

```

NOT plain voxel-index distance: this project's
voxels are
anisotropic (Z=1.0μm, XY=0.174μm typical), so an
unscaled EDT
would treat a "2 voxel" gap as 2.0μm along Z but
only 0.35μm
in-plane, a ~5.7x orientation-dependent
inconsistency
do_merge ← gap ≤ max_gap (μm)
if not do_merge: fall back to contact-area check
(dilate f by 1 voxel, count overlap with t) ≥
min_contact
if do_merge: absorb f into t, update t's cached vol/
centroid/bbox
if nothing merged this pass: stop
return masks, total_merges

```

min_contact's fallback branch is largely dead code in practice at any reasonable max_gap: the smallest possible EDT gap between two literally-touching-but-disjoint fragments is exactly one voxel spacing, always $\leq$ a positive max_gap — confirmed both by this structural argument and empirically against real production GT data (see the project's session history).

large_contact_merge(masks, large_contact) — a separate, later, size-agnostic merge: any two objects (regardless of size) sharing a contact area $\geq$ large_contact voxels get merged, catching blobs split through a thick junction rather than a thin neck. Same repeat-until-stable structure as safe-merge, contact measured the same way (1-voxel dilation + overlap count); the larger of the two objects' ID always survives.

final_min_size_cleanup(masks, gt_min, fraction=0.618) — the last-resort safety net, run after GMM/safe-merge/large-contact: removes anything still below $\text{fraction} \times \text{gt_min}$ (golden ratio by default — strict enough to reject debris, lenient enough not to reject a legitimately smaller-than-average real cell). Nothing upstream is guaranteed to remove every possible debris object (GMM judges by population statistics; safe-merge/large-contact only act when a mergeable neighbor exists), so this stage is a floor under all of them, not a replacement.

relabel_sequential(masks) — renumbers surviving labels 1..N with no gaps.

run_full_pipeline(volume, model_path, cellprob, flow, anisotropy, max_gap, min_contact, large_contact, gt_min, ..., flow_cache_path)
— glues every step above together:

```
1. dP, cellprob_map, shape ← predict_flows(...) or reuse
   precomputed_flows /
   a matching on-disk flow_cache_path (see save/load_flow_cache
   above)
2. masks ← masks_from_flows(...) #
   raw prediction
   raw_masks ← masks.copy() # kept in stats['raw_masks'] – this
   is what
                               # _widget.py always auto-saves as
   <stem>_cp_masks.tif
3. masks, gmm_cutoff, gmm_removed ← gmm_cleanup(masks)
4. masks, safe_merges ← krendl_safe_merge(masks, max_gap,
   min_contact, gt_min, scale_zyx)
   # this is what _widget.py always auto-
   saves as
   # <stem>_cp_masks_corrected.tif
5. masks, lc_merges ← large_contact_merge(masks, large_contact)
6. masks, final_removed ← final_min_size_cleanup(masks, gt_min,
   final_min_fraction)
7. masks, n_final ← relabel_sequential(masks)
return masks, stats{n_row, n_after_gmm, n_after_safe_merge,
n_after_large_contact,
                    n_after_final_min_size, n_final,
gmm_cutoff_vox, gmm_removed,
                    safe_merges, large_contact_merges,
final_min_threshold_vox,
                    final_min_removed, raw_masks}
```

rerun_single_cell(volume, labels, label_id, model_path, ..., pad_z=15, pad_xy=40) — “fix one mis-segmented cell” without a full-fish re-run: crops to that label’s own padded bounding box, re-runs the *entire* run_full_pipeline() on just the crop (safe to reuse unmodified — GMM cleanup already no-ops below 3 objects, and the merge/cleanup stages use a fixed gt_min floor, not crop-local population statistics), then splices back only the crop-result objects that actually overlap the original label’s own footprint — anything else the crop’s own do_3D pass happens to also detect (e.g. a neighboring cell caught by the padding) is discarded, never duplicated into the full volume. If the


```

return new_labels, report{sigma_xy, sigma_z, n_cells_total,
n_cells_sanded,
                                skipped_cells,
n_debris_fragments_removed}

```

Deliberately kept as its own small default (0.7/0.7 voxels), much smaller than the Pixel Classifier’s own Smooth σ_{XY}/σ_Z (1.5/3.0) — that pair decides whether raw blobs *merge* into one 3D instance before labels even exist; sanding runs after labels already exist and is foreign-protected (structurally can’t merge two cells), so its only job is polishing edges, not reshaping.

`_widget.py` — Tab 2 handlers

Route auto-detection: the Pixel Classifier and Cellpose-SAM sections (plus downstream tools — Resort/Split/Save) show or hide based on whether a `_NoBG` or `_ExtRm`-suffixed layer exists anywhere in the viewer, not just whichever layer is currently active — both can be shown together (e.g. comparing routes on the same fish).

`_on_create_labels()` (Pixel Classifier): background-thread wrapper reading Smooth σ_{XY}/σ_Z , Min volume, Min hole size sliders, calling `create_labels()` (3D Gaussian smooth → re-threshold → per-slice hole-fill → true 3D connected components → volume filter → renumber — no per-slice overlap-linking step), adding the result as a new `Labels` layer named `<image>_labels`.

`_on_run_cellpose_seg()` (Cellpose-SAM Segmentation): validates the checkpoint file (`_is_valid_cellpose_checkpoint()` — checks it’s a zip archive, the format `torch.save()` uses, to catch a wrongly-selected non-checkpoint file before a multi-hour run rather than deep inside `torch.load()`), auto-selects the matching `<stem>_brain_mask` layer if the “Auto-correct” checkbox needs skin protection, then runs `run_full_pipeline()` in a background thread wrapped in `capture_live_output` (surfaces Cellpose’s own otherwise-silent flow-following/hole-fill/small-object-removal step into the GUI log). On completion: adds the result as a `Labels` layer, **unconditionally** (no checkbox) writes `<stem>_cp_masks.tif` (raw_masks) and `<stem>_cp_masks_corrected.tif` (post-safe-merge) alongside the fish’s other outputs, and — if “Auto-correct labels via contrast sweep after segmentation” is ticked — chains into the Stage 3 auto-correction pipeline (see §5) as a further background stage.

`_on_rerun_single_cell()`: resolves the correct source Image layer and target Labels layer via `_resolve_rerun_layers()` (selected Labels layer wins if any name; else `<image>_cellpose_labels` → `<image>_labels` → a single same-shape Labels layer; ambiguity is reported, never guessed — fixes a real earlier bug where a labels layer not literally named `<stem>_cellpose_labels` was rejected), then calls `rerun_single_cell()` in a background thread and splices the result back into the live layer.

Sanding UI: one shared checkbox + Sigma XY/Z pair (Common Settings) drives sanding at every call site that regenerates a label's shape (Correct Label 2D/3D, Correct Adjacent Labels, and — chained as its own stage — the Cellpose-SAM auto-correct pipeline), not a separate control per tool.

5. Stage 3 — Label editing & correction

UI: Tab 3 — Edit MG Labels, working against a shared explicit Signal/Labels/Brain-mask layer selection (`_edit_labels_layer()`, see §1 “Layer-selection philosophy”) plus a Label A / Label B selector pair most tools reuse. **Modules:** `_labeling.py` (the editing primitives), `_grow_correct.py` (auto-grow/until-stable orchestration), `_contrast_sweep.py` (calibrating the intensity threshold these tools use), `_auto_correction.py` (the full unattended per-cell pipeline chained after Cellpose-SAM Segmentation). This is the largest, most-iterated part of the plugin.

Shared building blocks

Every tool in this section that regenerates a shape from intensity shares one convention: **candidate = image >= lo** — a one-sided threshold, not a band. `hi` (the signal layer's own display ceiling) is shown for context only, never used to restrict the mask — a narrow contrast window like `[100, 101]`, chosen to make the display saturate into a clean silhouette, would otherwise exclude the cell's true bright interior (well above `hi`, just shown as saturated white) if `hi` bounded the mask.

Every tool also shares **foreign-label protection**: any pixel already claimed by a *different* existing label is excluded from the candidate mask before anything (connected components, watershed) runs — a

foreign label can never be grown into, merged into, or even appear in a corrected shape, regardless of how the threshold or connectivity falls.

`_labeling.py` — editing primitives

`resort_labels(labels, sort_by, reverse)`: renumber every label 1..N by a chosen key — "size" (voxel count via `bincount`, natural = descending, largest→1), "centroid_z/y/x" (via `label_centroids_zyx`, natural = ascending), or "complexity" (skeleton branch count via `_statistics._skeleton_stats`, natural = descending, computed per-label in a thread pool). Only positive labels are ever touched by the final LUT remap — a sentinel like skin's -1 would otherwise corrupt via numpy's negative-index wraparound (`lut[-1]` = last element).

`label_centroids_zyx(labels, label_ids)`: bit-identical to `scipy.ndimage.center_of_mass`, but computed from each label's own bounding box (`find_objects` + a per-label crop) instead of one whole-volume call — ~10-100x faster when a large negative sentinel (skin) is present, since scipy's own implementation takes a much slower path for negative-containing arrays.

`remove_debris(labels, threshold, skin_label_id=None)`: zeros every *spatially connected fragment* smaller than `threshold` voxels — evaluated per connected component (26-connectivity), not per raw label ID, so several small disconnected fragments still sharing one old ID (e.g. after a manual edit) are each judged on their own size, not their summed total. Does **not** renumber survivors. `skin_label_id`, when given and present, is swept too (skin's own bounding box computed manually since `find_objects()` never indexes negative values). The per-label sweep (`_remove_debris_from_crop`) clears every small fragment of a label in one `bincount`+boolean-index pass, not one array scan per fragment — a real historical bug (per-fragment `crop[cc==id]=0` scanning the whole crop each time) made this "hang" on a fragmented skin label with thousands of small pieces.

`split_label(labels, target_label, n_splits, sigma, min_distance, mode, z, image)` — watershed split into N parts:

`mode="3d"`: surface = the mask's own EDT (splits at the geometrically

 narrowest neck, regardless of what the underlying signal looks like)

`mode="2d"`: surface = raw signal intensity on slice `z` (seeds at brightness

peaks; cut follows the dimmest ridge – for when two things only

merge on ONE cross-section, where the 3D mask has no real neck

but the signal already dips)

1. crop to bbox(+pad)
2. Gaussian-smooth the surface (GPU if available)
3. `h_maxima(surface, h)` – TOPOLOGICAL prominence, not `peak_local_max`'s

Euclidean-radius search: finds peaks separated by the thinnest saddle

regardless of how physically close their centers are. Auto-reduces `h`

from 50% of max down to 0.5% until `>= n_splits` peaks are found.

4. within each `h_maxima` region, seed = its own EDT/intensity maximum;
- drop seeds closer than `min_distance` to an already-chosen one
5. `watershed(-surface, markers=seeds, mask=mask_crop)`
6. clear the 1-voxel interface where two parts directly touch (outer surface of each part untouched)

Part 1 keeps `target_label`; parts 2..N get fresh IDs above the current max.

`join_labels(labels, label_a, label_b)`: the inverse — every `label_b` voxel becomes `label_a`. One vectorized boolean assignment, no crop/GPU needed.

`correct_label_from_intensity(labels, image, label_id, z, lo, hi, pad)` — Correct Label's 2D engine: within `label_id`'s own padded bbox on slice `z`, `candidate = image >= lo`, foreign-excluded, painted in wholesale (not connectivity-restricted to the label's existing footprint — a human reading the same contrast window would count a visible blob as part of the cell whether or not it happens to already touch). Raises rather than silently emptying the label if nothing survives.

`correct_label_from_intensity_3d(labels, image, label_id, lo, pad, ..., auto_grow, growth_step, max_iterations, until_stable, max_stability_passes, resolve_adjacent, z_extent_pad)` — Correct Label's 3D engine, the most complex function in the plugin. NOT a bounding box over the whole cell — loops one Z-slice at a time, each with its own locally re-derived neighborhood:

```
zmid ← z_orig_min + ceil((z_orig_max - z_orig_min) / 2)
walk z from zmid down to z_orig_min, then from zmid+1 up to
```

```

z_orig_max:
    _or_seed(z): # before correcting z, OR-union z-1 and z+1's
footprints
                # (within the label's fixed original range) into
z's own
                # current footprint as watershed seed context –
an
                # already-walked neighbor contributes its JUST-
corrected
                # shape, a not-yet-walked one its ORIGINAL
shape. OR, not
                # AND (the source pseudocode's own literal "AND"
was a
                # stated user error – AND would erase real
signal any
                # time a neighbor's shape isn't already
identical).
                # Pasted only into empty background or skin
(-1), never
                # a different real label's territory.
    _correct_slice_full(z):
        if another real label is present in z's own local area:
            jointly re-derive via the same marker-seeded watershed
            Correct Adjacent Labels uses
(_correct_label_group_2d_core,
            focus_ids=[label_id]) – only label_id's OWN resulting
side
            is ever written back
        else:
            plain single-label 2D correction
(_intensity_correct_2d)
            if auto_grow and result touches its own local window's
edge:
                retry this ONE slice at pad += growth_step, up to
max_iterations
            if until_stable:
                keep re-running THIS SLICE (reseeded from its own last
result)
                until it matches the previous attempt, or
max_stability_passes
            extend outward past the original [zmin, zmax] one slice at a time
            (up to
            z_extent_pad, a FIXED value never inflated by growth): copy
            the

```

current footprint forward (foreign-protected), then run the same

local correction there; undo and stop the moment nothing supports it

– the true edge has been found

`remove_debris_for_label()` cleans up any small disconnected leftover

Every slice's own correction is *locally anchored* — re-derived fresh from wherever the label actually sits one slice ago, so it can only ever reach pad beyond its own verified position, never balloon into unrelated distant signal the way one whole-cell-sized window would. `resolve_adjacent=False` treats every other label as purely excluded territory instead of jointly resolving against it — used for skin's own sentinel, which borders essentially every real cell and shouldn't be watershed-split against each one.

`sand_label(labels, label_id, sigma_xy, sigma_z, pad)` — the per-label engine behind §4's sanding pass: 3D Gaussian-blur the label's own binary mask, re-threshold at 0.5, foreign-excluded; skips (reports why) if the result collapses to nothing or loses contact with the original footprint.

`_intensity_correct_2d / _intensity_grow_2d` — the shared single-slice threshold core every 2D tool above calls into: crop to `seed_mask's` `bbox+pad`, `candidate = image >= lo` minus foreign territory, fill the *whole* crop (not connectivity-restricted to the seed).

`correct_adjacent_labels_2d(labels, image, label_a, label_b, z, lo, pad, sigma)` — corrects two touching labels on one slice *simultaneously*, boundary placed by a **marker-seeded** watershed (markers = each label's own existing footprint, not auto-detected peaks — Split Label's blind peak-finding was tried and rejected here, since a combined region can have more than one real intensity dip and lock onto the wrong one). The working rectangle is sized from `label_a` alone (+pad), not the union — `label_b` only participates via whatever falls inside that rectangle.

`correct_label_group_2d / _correct_label_group_2d_core` — the N-label generalization of the above, used by Correct Adjacent Labels' own group-correction and (crop-only, no full-volume copy) by `correct_label_from_intensity_3d's` per-slice joint branch. `focus_ids` optionally scopes the working rectangle to a subset of the group (mirrors `label_a`-only scoping); markers are always each label's own

existing footprint, watershed runs unmasked over the whole crop then intersects with the real candidate signal so a diagonally-touching island still gets an owner.

touching_groups_for_stack (general-purpose utility, not on the current pipeline's own critical path — replaced there by the per-cell engine's own adjacency resolution): per-Z-slice union-find of mutually-touching label clusters (`_touching_pairs_on_slice + _union_find_groups`).

copy_label_to_adjacent_slice(labels, label_id, z_src, direction): copies one label's 2D footprint onto the next/previous slice, clearing its own old footprint there first, foreign-excluded (reports how many pixels were dropped to foreign territory rather than pretending an exact copy).

Skin protection: `skin_voxel_count(labels, skin_label_id=-1)` — the single, array-content-only definition of “is skin already protected” every path checks (never widget state). `seed_skin_label(labels, brain_mask, skin_label_id=-1)` — bulk-fills every background voxel outside the brain mask with the sentinel -1 (never overwrites a real label, never touches inside the brain); this is only a coarse seed. `trim_skin_label(...)` — a thin wrapper around `correct_label_2d_stack` (below) trimming that bulk seed down to real signal-supported territory, per-slice only (skin has no cross-slice continuity problem to solve and spans nearly the whole Z range, so the 3D walk's machinery would be pure waste — the same reason 3D correction of an *existing* skin label is structurally blocked in the interactive UI). No brain-mask clamp — a legitimate inward correction along skin's real inner boundary is never blocked; a `remove_debris(..., skin_label_id=...)` pass run right after sweeps up any small stray blob skin absorbed instead. `remove_label(labels, label_id)` — plain bulk clear, not connectivity-restricted, used to drop skin's sentinel when no longer needed.

correct_label_2d_stack(labels, image, label_id, lo, pad, ..., n_workers) — corrects one label across its *whole* Z range, slice by slice, **independently and in parallel** (ThreadPoolExecutor, 75% of cores by default, `n_workers` overridable so a caller already parallel across cells — the auto-correct waves — doesn't multiply thread pools), each slice via `grow_correct_label_2d()` (the exact function the interactive Correct Label 2D button calls) on a one-slice *view*, not a full-volume copy per slice — keeping RAM bounded regardless of how many slice threads run concurrently. Wherever another real label is present, the boundary is jointly resolved via the same marker-seeded watershed, but only `label_id`'s own resulting side is ever painted back —

guaranteed by watershed's own marker semantics (a label's own existing pixels are always its own marker, which watershed can never reassign), not a separately-enforced rule. `trim_skin_label()` is a thin wrapper calling this with `label_id = skin_label_id`.

`_grow_correct.py` — auto-grow/until-stable orchestration

`grow_correct_label_2d(labels, image, label_ids, z, lo, initial_pad, growth_step, max_iterations, sigma, until_stable, max_stability_passes)` — wraps `correct_label_group_2d` in two nested loops:

```
group ← {label_ids}  (a single int for Correct Label; [A, B] for Correct
                        Adjacent Labels – the rectangle stays
scoped to
                        the FIRST id even as the group grows)
for each stability pass (1 if until_stable=False, else up to
max_stability_passes):
    for attempt in 1..max_iterations:
        correct the current group at the current pad
        new_neighbors ← labels the ORIGINAL group's own footprint
(not a
                        folded-in neighbor's) GENUINELY touches,
that
                        aren't in the group yet
        if new_neighbors: group |= new_neighbors; redo at the SAME
pad
        elif the focus id(s) still touch their own rectangle's
edge:
            pad += growth_step; retry
        else: converged, break
    if until_stable and this pass's result on every group member
matches
        the immediately preceding pass's: stop (stable)
    else: reseed the next pass from this pass's own result
```

Neighbor discovery is deliberately narrow: only the *originally requested* label's own genuine touching-adjacency counts (never “merely present in the box,” which would keep finding something as the box grows; never an already-folded-in neighbor's own further touches, which could cascade the group outward indefinitely).

grow_correct_label_3d(labels, image, label_ids, lo, ...) — a thin pass-through to **correct_label_from_intensity_3d(auto_grow=True, ...)**: all growth and stability happens *inside* that function's own per-slice walk, never as a whole-cell retry at a bigger global pad.

format_grow_report(report, mode): renders either shape into the same human-readable style (which slices grew, which took multiple stability passes, whether it converged) used throughout the plugin's status/report boxes.

_contrast_sweep.py — calibrating Correct Label's own lo

Answers “which lo best *reproduces* what Cellpose-SAM already segmented” — self-referential calibration against the model's own output, not independent GT (a deliberately different target from every GT-sweep tool in §8).

select_calibration_samples(labels, scale_zyx, n_cells, slices_per_cell, edge_margin_um): picks the `n_cells` most morphologically complex cells (skeleton branch count) whose centroid sits $\geq$ `edge_margin_um` from the volume's own edge (a proxy for “not already touching skin”), each axis's margin capped at 30% of that axis's own extent so an un-capped physical- μm margin doesn't collapse Z's own thin window to nothing. For each, draws `slices_per_cell` Z-slices spread across the middle 60% of its own Z-extent (avoiding thin top/bottom cross-sections).

default_lo_candidates(image, samples, pad, n_steps): auto-scales the candidate range from the 1st–99.5th percentile of intensity *around the actual samples*, not a hardcoded guess — robust to a different channel/bit-depth/normalization.

sweep_contrast_lower_value(labels, image, samples, lo_candidates, pad):

```
for each candidate lo:
    for each (label_id, z) sample:
        corrected ← _intensity_correct_2d(..., lo) result (IoU=0
on error,
                                doesn't abort the whole sweep)
        iou[sample] ← IoU(corrected, the sample's EXISTING
footprint)
    mean_iou[lo] ← mean(iou over all samples)    # jointly, not
each sample's
```

```

# own
independent best –

# stops one
outlier sample

# from swinging
the result
best_lo ← argmax(mean_iou)

```

Any non-positive label (skin -1) is treated as plain background during this sweep — otherwise the correction-under-test would see skin as foreign territory and exclude it, making the calibrated lo depend on whether skin happened to already be protected.

`_auto_correction.py` — the full unattended per-cell pipeline

`auto_contrast_correct_stack(labels, image, scale_zyx, brain_mask, skin_image, ...)` — chained onto the end of Cellpose-SAM Segmentation (§4) when its “Auto-correct” checkbox is ticked, or run standalone via “Auto-correct Existing Labels” against any Labels layer. Six steps:

```

1. best_lo ← sweep_contrast_lower_value(...) via
   select_calibration_samples/
       default_lo_candidates (or lo_override, skipping the
sweep
       entirely, minus an optional lo_adjustment subtracted
after)

2. if skin (-1) already present in `labels`: REUSE it exactly as-
is,
       never re-seed/re-trim (protects any prior manual
correction)
   else:
       seeded ← seed_skin_label(labels, brain_mask)
       labels_with_skin, skin_report ← trim_skin_label(seeded,
skin_image
       or image, -1, best_lo, pad=skin_pad, ...)
       # same best_lo real cells get -- skin's own joint-
resolution
       # against a touching cell (§ above) needs no separate
margin

3. remove_debris(labels_with_skin, threshold, skin_label_id=-1)

```

```

# sweeps whatever small stray blob skin's own unclamped trim
absorbed

GATE: skin_voxel_count(...) == 0 after steps 2-3 → raise. No
cell is
EVER corrected without verifiably-present skin protection.

4. resort_labels(sort_by="centroid_z")
   touching_skin_ids ← _labels_touching_skin(...) # GEOMETRIC
test on
   skin's FINAL territory (dilate each cell by skin_touch_px,
check
   overlap) – NOT read from trim_skin_label()'s own
foreign_nearby
   report, which lists every label ever transiently folded
into a
   growth attempt (found to wrongly flag 33/33 cells on a real
fish)
   waves, boxes ← _compute_correction_waves(...) # partitions
every
   cell into provably non-conflicting groups by each cell's
own
   worst-case reach (pad + growth_step*max_iterations); same-
wave
   cells run concurrently (75%-of-cores ThreadPoolExecutor),
later
   waves see all earlier waves' already-corrected state

5. for each wave, for each cell (in parallel):
   if touching skin: correct_label_2d_stack(...,
label_id=cell)
                                   # same per-slice mechanism
trim_skin_label()
                                   # itself uses – a cell meeting skin has
to do
                                   # so on skin's own 2D-only terms
   else:
       grow_correct_label_3d(..., cell)
       merge only this cell's own expanded box back into the
shared array
       (wave members' boxes never overlap, so this is race-free)

6. remove_debris(new_labels, threshold, skin_label_id=-1) #
final
   whole-layer safety net

```

(report-only) `_flag_possible_non_microglia(...)`: flags any cell with
 `>= non_microglia_fraction` (default 40%) of its volume
outside the
 brain mask, OR of its own boundary SURFACE within
skin_touch_px of
 skin – macrophages or MONAI-missed skin that Cellpose
segmented as
 a cell. Nothing is removed or changed; report only.

`format_auto_correction_report()` renders one consolidated report: the calibration, skin protection outcome, wave/parallelism summary, per-cell detail (reusing `format_grow_report()` for 3D cells, a matching `_format_2d_vs_skin_report()` for skin-touching ones), and the non-microglia flag list.

`_widget.py` — Tab 3 (Edit MG Labels) handlers

Shared layer selection: `_edit_labels_layer()` / `_edit_signal_combo / _edit_mask_combo` — one explicit combo triple the user sets once, replacing the old implicit “whatever’s active” fallback (see §1).

`_on_resort_labels()`: reads sort-by/reverse from the UI, calls `resort_labels()` in a background thread, writes the result back in place (`layer.data[:] = result`; `layer.refresh()` — not `layer.data = result`, which was found to race with napari’s own async thumbnail update and crash; see §9).

`_on_remove_debris()`: reads Min volume/Min hole size from Common Settings, calls `remove_debris()`.

`_on_split_label()`: reads target label, N splits, σ , min distance, mode (2D/3D), calls `split_label()`, adds the new IDs to the layer.

`_on_join_labels()`: reads Label A/B (each with a “Use selected” button reading `layer.selected_label`), calls `join_labels()`.

`_on_correct_label()`: reads the signal layer’s live `contrast_limits` for `lo`, target label, pad, mode (2D uses the current slice; 3D disabled outright for label -1, since a 3D correction of skin would walk nearly the whole fish — see `correct_label_from_intensity_3d`’s own note), auto-grow/until-stable checkboxes; calls `grow_correct_label_2d/3d` accordingly; sands the result afterward if Sanding is enabled (§4).

`_on_correct_adjacent_labels()`: reads Label A/B + current slice, calls `grow_correct_label_2d(label_ids=[A, B], ...)`, sands both labels afterward if enabled.

`_on_copy_label_to_adjacent_slice()`: reads label + direction, calls `copy_label_to_adjacent_slice()`.

`_on_protect_skin()`: runs the contrast sweep (§ above, “Calibrate Correct-Label Contrast” is the canonical source — every other autosweep tool in the plugin reuses its result rather than re-sweeping) to get `best_lo`, then `seed_skin_label + trim_skin_label`; refuses (“Skin is ALREADY protected...”) if `skin_voxel_count() > 0` already. Populates the “Reduce `best_lo` by” reduction slider’s own section.

`_on_remove_skin_label()`: calls `remove_label(labels, -1)`.

`_on_auto_correct_existing_labels()`: background-thread wrapper around `auto_contrast_correct_stack()` against whatever Labels/Image/brain-mask layers are selected — the same pipeline Cellpose-SAM Segmentation’s own “Auto-correct” checkbox chains into, exposed as its own standalone button so it can be re-run against any existing Labels layer, not just fresh segmentation output. Wrapped in the `hang watchdog (_hang_watch_start/stop, 30-minute)` `faulthandler.dump_traceback_later()` given how long a full-fish run can take.

`_on_sand_label()` (interactive single-label sanding, when not chained automatically): reads σ_{XY}/σ_Z from the shared Sanding section, calls `sand_label()`.

“Hide skin label” toggle: a purely visual `DirectLabelColormap` swap (never touches `layer.data`) making skin render transparent — a `_FastDirectLabelColormap` subclass bypasses napari’s own slower numba-dict backend selection for this one colormap (a real ~2x measured fix for 3D-display slowness), verified to produce byte-identical output to the base class.

“Drift View in 3D”: Start/Stop + speed slider driving a `QTimer`-ticked camera rotation (`viewer.camera.angles`, composed via a real `scipy.Rotation` each tick rather than incrementing 3 Euler angles independently — napari’s own camera API restricts its middle axis to $\pm 90^\circ$, which an increment-based approach visibly stalls against). Camera-only; never touches any layer’s data.

6. Stage 4 — Statistics

UI: Tab 4 — Statistics, with a scrollable per-column output checklist (defaults matching core microglia morphology; several columns Nathalie's review flagged as marginal/unvalidated default OFF — see §9). **Module:** `_statistics.py`, one function (`compute_stats`) producing a `pandas.DataFrame`, one row per label, up to 51 columns. **Goal:** turn a corrected Labels layer into a per-cell CSV.

Three-phase pipeline (`compute_stats`)

Phase 1 — batch regionprops, GPU or CPU (`_batch_regionprops`): one vectorized pass over every label at once —

`cucim.skimage.measure.regionprops_table` on GPU (CuPy) if available, else `skimage.measure.regionprops_table` on CPU; both requesting the same property list (`label`, `area`, `centroid`, `bbox`, `inertia_tensor`, `inertia_tensor_eigvals`, `axis_major_length`, `axis_minor_length`, `solidity`, `extent`). From this table, vectorized post-processing derives:

```
IT ← the 3x3 inertia tensor per label, stacked
eigvals, eigvecs ← eigh(IT)                # symmetric
eigensolve, all labels at once
principal_axis_dir[i] ← "Z"/"Y"/"X", whichever axis the LONGEST
eigenvector                                (column 0) has its largest |component|
along
axis1_um ← axis_major_length × (that axis's own physical scale)
axis3_um ← axis_minor_length × mean(sz, sy, sx)
axis2_um ← (axis1_um + axis3_um) / 2        # not
independently measured
elongation ← axis1_um / axis3_um            # 1 = sphere, >1 =
elongated
bbox_dz/dy/dx_um ← bbox extent × physical scale
volume_um3 ← voxel count × (sz × sy × sx)
eq_diam_um ← (6 × volume_um3 / pi) ^ (1/3)  # equivalent-
sphere diameter
```

Phase 2a — per-label marching cubes + skeleton, threaded

(`_slow_stats_worker`, one call per label in a `ThreadPoolExecutor` at 50% of cores): for each label's own bbox crop,

```
_surface_area(binary, scale_zyx):
    pad the crop by 1 voxel (avoids a boundary artifact)
    verts, faces ← marching_cubes(padded, level=0.5,
```

```

spacing=scale_zyx)
    return mesh_surface_area(verts, faces)          #  $\mu\text{m}^2$ , 0.0 on
any failure

_skeleton_stats(binary, scale_zyx):
    skeleton ← skeletonize(binary)                  # skimage,
topological thinning
    sk ← skan.Skeleton(skeleton, spacing=scale_zyx)
    bd ← skan.summarize(sk)                          # one row per
branch segment
    n_branches ← len(bd)
    n_endpoints ← count of branch-type == 1 (free ends, not
junction-junction)
    mean/max_branch_len_um ← mean/max of bd's euclidean-distance
column
    tortuosity ← mean(path_length / euclidean_distance) over
branches with
                                nonzero euclidean distance – 1.0 = perfectly
straight
    returns (0,0,0.0,0.0,1.0) on any failure (e.g. skan not
installed, empty skeleton)

```

Phase 2b — intensity stats, optional, also threaded

(`_intensity_stats_worker`, only if an Image layer was given): per-label mean/sum/coefficient-of-variation of the raw signal intensity within the label's own mask.

Phase 3 — assembly + derived metrics + descriptions:

```

for each label:
    sphericity ← min(1.0, pi^(1/3) x (6 x volume)^(2/3) /
surface_area)
    surface_to_volume_ratio ← surface_area / volume
    branch_density ← n_branches / volume x 1e6          # per- $\mu\text{m}^3$ ,
scaled for readability
    endpoint_density ← n_endpoints / volume x 1e6
    process_complexity ← n_endpoints x mean_branch_len / volume
    morphotype ← _classify_morphotype(sphericity, solidity,
elongation, n_branches, sav)
    description ← desc_fn(row)                          # rule-
based or an LLM backend

```

_classify_morphotype(spher, solid, elong, n_br, sav) — rule-based, priority order:

```

if elong > 3.5 and n_br <= 3:                                → "Rod-
shaped"
elif spher > 0.70 and solid > 0.80 and n_br <= 2:          →
"Amoeboid"          (activated)
elif n_br >= 6 and sav > 2.0 and spher < 0.55:              →
"Ramified"          (resting)
elif n_br >= 4 and spher < 0.65:                             →
"Intermediate-ramified"
else:                                                         →
"Intermediate"

```

Post-assembly, whole-dataset spatial statistics (`_spatial_stats`,
`scipy.spatial.cKDTree` over all centroids at once):

```

tree ← cKDTree(centroids_um)
k=3 query → self + 1st NN + 2nd NN (k=2 fallback if only 2 cells
exist)
nearest_neighbor_dist_um, nearest_neighbor_2_dist_um ← those
distances
Clark-Evans 3D index:
    density ← N / bbox_volume(all centroids)
    expected_NND ← Gamma(4/3) x (3 / (4*pi*density))^(1/3)    #
Poisson-process expectation
    nearest_neighbor_ratio ← observed_NND / expected_NND      # <1
clustered, >1 dispersed
local_density_100um ← count of other centroids within a 100µm
radius (tree.query_ball_point)
depth_normalized ← centroid_z_um / max(centroid_z_um), clipped to
[0,1]

```

Post-assembly, optional brain-region assignment

(`_assign_brain_regions` + `_polyline_side_and_dist`), only if a Shapes
layer with "line"/"path" shapes and region names were given:

```

boundaries ← the drawn polylines, in µm, sorted by mean X of their
own vertices
                (ascending = most anterior first) -- the fish lies
along X,
                head at X=0; each boundary is meant to run top-to-
bottom (Y),
                dividing anterior (smaller X, left of the path) from
posterior
for each cell's centroid:
    for each boundary, nearest-to-first:
        find the closest point on ANY of the boundary's segments

```

```

(point-to-
    segment projection, not just to its two endpoints --
supports a
    curved multi-vertex polyline, not just a straight 2-point
line)
    side ← sign of the cross-product at that nearest segment
    if side says "posterior of this boundary": region_index +=
1
    region ← names[region_index]                # N boundaries
-> N+1 regions
    region_boundary_dist_um ← distance to the closest boundary
overall

```

Description backends (`_make_desc_fn`)

Four interchangeable backends, selected by `backend_config["backend"]`: - **rule** (default, always available, offline): `_rule_based_description` — a template assembling shape (from sphericity/elongation), surface texture (from solidity), branching, and the morphotype tag into one sentence. - **ollama**: POSTs the same statistics, formatted into `_STATS_PROMPT`, to a local Ollama server's `/api/generate`. - **openai** / **claude**: same prompt to the OpenAI Chat Completions or Anthropic Messages API respectively, using a key from `_secrets.py` (§2). All three network backends catch their own request errors and return an inline "[Backend error: ...]" string rather than raising — one label's description failure never aborts the whole run.

`_widget.py` — Tab 4 handlers

Per-column checkboxes: `_COL_GROUPS` maps one checkbox to multiple CSV columns where they're conceptually one unit (e.g. `centroid_vox` → all three `centroid_z/y/x_vox` columns; `bbox_vox` → all six `bbox` corner columns; `nn_1st/nn_2nd` → the label+distance pair each). Validation warns (doesn't crash) if brain-region or intensity columns are checked but no Shapes/Image layer is selected, and just skips those columns.

`_refresh_stats_layers()`: repopulates the Image/Shapes layer combos on every napari layer add/remove event, preserving the current selection across the refresh (`blockSignals()` guard).

`_on_generate_stats()`: reads the active Labels layer's own scale from `self._state["metadata"]` first (the value Tab 1's Open button set), falling back to the Labels layer's own napari scale only if no stack metadata is on record — a real historical bug had this backwards

(defaulting to the Labels layer's own possibly-unset scale), silently making `centroid_um/volume_um3` identical to their voxel-unit counterparts whenever a `labels.tif` was loaded without its own embedded scale. Extracts region-boundary polylines from the selected Shapes layer via `_extract_region_lines_um()` (accepts both "line" and multi-vertex "path" shape types), calls `compute_stats()` in a background thread, filters the returned DataFrame down to only the checked columns before writing the CSV.

GT-verification checkbox ("This is verified ground truth"): when ticked, also measures the active layer's own smallest/largest real cell volume and smallest real hole directly from this GT, updating the shared cross-fish Min/Max-volume and Min-hole-size histories (`_update_gt_history`) other tools throughout the plugin read from — the same one-shot, self-unticking pattern every GT-sweep tool in §5/§8 uses.

7. Stage 5 — AI Tools

UI: Tab 5 — AI Tools, a manual MONAI-Training/Cellpose-SAM-Training switch (not auto-detected — no per-layer signal to key off here), gated informationally (not blocked) on GPU/VRAM (§2's `_gpu_check.py`). **Goal:** build and train the two models Stages 1–2 depend on, and produce the ground-truth data that trains them.

`_gt_annotation.py` — polygon-based GT annotation

Plain viewer-taking functions on the plugin's own shared viewer, rather than a second `napari.Viewer()`: the target Image layer is passed explicitly (not "first Image layer in the list," a real risk once Tabs 1–3 have added other layers); the `brain_polygons` Shapes layer is auto-created on demand; `generate_masks()` leaves every other layer in the viewer untouched.

```
resample_polygon_preserve_order(pts, n_points=96):
    walk the closed polygon's own perimeter by cumulative arc
length,
    resample to exactly 96 evenly-arc-spaced points – gives every
hand-drawn polygon, regardless of how many vertices were
clicked,
    the same point COUNT and roughly consistent point
CORRESPONDENCE
    for the interpolation step below
```

```

is_polygon_clockwise(yx) / standardize_polygon_direction(...):
    shoelace-formula signed area – every key-slice polygon is
forced to
    the SAME winding direction as the first (reference) one, since
    point-to-point interpolation between two oppositely-wound
polygons
    would twist through the middle instead of sweeping smoothly

interpolate_shapes(viewer, image_layer):
    by_z ← {slice: resampled 96-point polygon} for every hand-
drawn key polygon
    standardize every polygon's winding to match by_z's own first
slice
    for every Z slice in the volume:
        if a key polygon exists exactly here: use it as-is
        elif slice > PROPAGATION_SLICE (default 90) and
propagation is on:
            reuse the PROPAGATION_SLICE polygon UNCHANGED (just
relabel its
            own Z) – for the common case where the brain
silhouette stops
            changing meaningfully past a certain depth, so no
further
            hand-drawn key slices are needed there
        else:
            find the nearest key slice below and above this one
            linearly interpolate every one of the 96 point-pairs
between
            them by the fractional Z position (point i of the
lower
            polygon -> point i of the upper polygon)
        → a new "brain_polygons_interpolated" Shapes layer, one
polygon per slice

generate_masks(viewer, image_layer, input_path):
    brain_mask ← rasterize every interpolated polygon via napari's
own
        Shapes.to_labels(), thresholded > 0
    optional median_filter along Z (window 3) to smooth slice-to-
slice
        rasterization jitter
    skin_mask ← 1 - brain_mask
    brain_only / skin_only ← image x each mask

```

```
    save brain_mask/skin_mask/original/brain_only/skin_only .tif +
both
    polygon .npz files under <input>/<stem>/ (the plugin's own
output-folder convention)
```

_xyz_patches.py — 2D orientation-slice crop generation (current training methodology)

Generates the actual crop set every real Cellpose-SAM training run since May 2026 has used — **not** the older bbox-based single/double/triple/quadruple crop approach a separate, still-available tool ports from an April-2026 script that was abandoned after failing (catastrophic forgetting on too little data).

```
stretch_z(slice_2d, anisotropy): bilinear zoom of the Z axis by
    anisotropy – brings a Z-cross-section (XZ or YZ) to roughly
the same
    pixel scale as a native XY slice, since this project's voxels
are
```

```
    strongly anisotropic
stretch_z_mask(...): the same zoom, nearest-neighbor (order=0) –
integer
    labels must never be blended/interpolated
```

```
_random_crops(img_slice, gt_slice, crop_size, n_crops, rng,
min_gt_pixels):
    repeat (up to 30x n_crops attempts):
        pick a random GT-positive pixel as an anchor
        place a crop_size x crop_size window with a random offset
so the
        anchor lands somewhere inside it (not necessarily
centered)
        keep it only if it contains >= min_gt_pixels of real GT
signal,
        and its own top-left corner hasn't already been used
```

```
generate_xyz_patches(image_path, gt_path, out_dir, anisotropy,
crop_size=512,
                    ncrops_per_slice=5,
max_per_orientation=320, ...):
    for each orientation in (XY native, XZ stretched, YZ
stretched):
        shuffle every GT-containing slice of that orientation
        for each slice (until max_per_orientation total crops
```

```

collected):
    draw up to ncrops_per_slice random crops from it
    cap the COLLECTED list to max_per_orientation, THEN write
files
    (not an early-exit mid-slice that could overshoot the
cap)
    save <orientation>_<slice:03d/04d>_<i:02d>.tif +
_masks.tif pairs

```

_crop_truncation.py — cleaning incidental-neighbor truncation

A crop framed around one target cell often also grazes the corner of a different, nearby cell purely by chance, keeping a valid-looking but tiny sliver whose visible centroid is nowhere near that cell's true center — a genuinely wrong training target for Cellpose's flow-vector loss (which points every voxel toward its own object's true centroid). This is the exact fix applied by hand to this project's real training data before being ported into the plugin.

```

clean_crop_truncation(crop_dir, gt_full_path, anisotropy,
threshold=0.9):
    back up crop_dir first (skip if a backup already exists –
never
    clobbers a still-valid backup with already-modified files)
    for each <stem>.tif/<stem>_masks.tif pair (xy/xz/yz_NNN_NN
naming):
        true_slice ← the SAME stretch_z_mask() transform
generate_xxyz_
        patches() used at generation time, applied to the
fish's own
        full GT volume at this crop's source slice index
(cached –
        many crops share a source slice)
    for each label present in the crop's mask:
        visible ← pixel count in the crop
        true_count ← pixel count in the true full-slice cross-
section
        if visible < threshold x true_count: zero this label
in the
            crop (an incidental neighbor caught mid-
truncation) –
            the crop's own INTENDED target cell is essentially
never

```

affected, since it's already near-complete by construction
overwrite the mask file only if something was actually zeroed

`_branch_calibration.py` — measuring real branch radius from GT

Recalibrates `train_xyz.py`'s `branch_radius` parameter (the erosion-survival-distance threshold the branch-weighted training loss uses to decide “this pixel is part of a thin branch”) from actually-measured GT morphology instead of a frozen guess.

```
measure_branch_diameters(gt_labels, scale_zyx, pad, min_volume):
    for each labeled cell (cropped to its own padded bbox):
        edt ← anisotropic distance_transform_edt(binary_mask,
sampling=scale_zyx)
        skel ← skeletonize(binary_mask)                                #
topological thinning
        skel_valued ← edt VALUES baked onto the skeleton's own
nonzero
        voxels (a real skan API gotcha: Skeleton(...,
source_image=...))
        does NOT feed skan.summarize()'s own mean-pixel-value
column –
        that column reads the skeleton_image array's own
values, so
        the EDT has to be baked IN, not passed as a separate
argument)
        bd ← skan.summarize(Skeleton(skel_valued,
spacing=scale_zyx))
        diam_um ← 2 x bd's mean-pixel-value per branch segment
        → pooled diam_um / length_um across every cell

recommend_branch_radius(gt_labels, scale_zyx, xy_scale=None):
    stats ← measure_branch_diameters(...)
    tip_diam_um ← the thinnest QUARTILE of all measured diameters
        (distal branch tips – what branch_radius exists
to protect)
    tip_radius_um ← mean(tip_diam_um) / 2
    recommended_px ← round(tip_radius_um / xy_scale)    # μm ->
pixels at the
        crop's own XY resolution, directly
```

comparable to

`train_xyz.py`'s own `branch_radius` parameter

An approximation, not an exact reconstruction — the training-time threshold is a Chebyshev (chessboard) erosion distance in pixels; this measures a true Euclidean physical radius and converts. Close enough to calibrate a threshold, acknowledged in the docstring as not identical.

`_training_jobs.py` — cross-platform detached process management

Both MONAI (`train.py`, hours–days) and Cellpose-SAM (`train_xyz.py`, ~20h) training runs must survive napari closing and be re-discoverable after a full restart. Built on primitives that behave identically on Linux/Mac/**Windows** (this plugin ships a native Windows install path, so a tmux-only design was rejected): a detached subprocess. Popen + PID tracking via `psutil`, `stdout/stderr` to a plugin-controlled log file the GUI tails.

```
launch_detached(argv, cwd, log_path, conda_env, notify=None):
    if notify is given:
        write a small STANDALONE supervisor .py file (stdlib-only
        —
        subprocess/smtplib/re, no dependency on this package
        being
        importable) that: runs the real "conda run -n <env>
        <argv>"
        command (same log_path the GUI already tails live) ->
        waits
        for it to exit -> parses the log for the best-
        checkpoint
        metric -> emails a completion report via SMTP_SSL
        the SUPERVISOR itself is what's launched detached, not
        argv
        directly — so the email still fires even if napari
        closes
        and never reopens before the job finishes.
    kill_process_tree
        needs no special case: the real training process is a
        descendant of the supervisor, so killing the whole
        tree (a
        manual Stop click) kills the supervisor before it
        reaches
        its own email step, correctly sending no "stopped"
```

```

email.
    else:
        Popen(["conda", "run", "-n", env, "--no-capture-
output", *argv],
              stdout/stderr -> log_path,
              start_new_session=True on POSIX (setsid) /
              CREATE_NEW_PROCESS_GROUP|DETACHED_PROCESS on
Windows)
        return pid

is_running(pid) / kill_process_tree(pid): psutil-based liveness
check /
    recursive terminate-then-force-kill of pid + every descendant
(conda
    run spawns a child python process, so killing only the top PID
would
    leave training running)

tail_log(log_path, n_bytes=8192): last N bytes only, bounded
regardless
    of how many days/GB the log has grown to

patience_exceeded(log_path, metric_cfg, patience):
    parse every (epoch, metric_value) the log's own regex matches
    (MONAI_METRIC: "Epoch N Summary:" ... "Full-brain Dice: X
    [MODEL SELECTION]", DOTALL, higher-is-better;
CELLPOSE_METRIC:
    "N, train_loss=..., test_loss=X" on one line, lower-is-
better –
    the ONLY thing that differs between the two scripts'
early-
    stopping is this metric config, the counting logic is
shared)
    best ← the best-value checkpoint logged so far
    checkpoints_since_best ← how many logged checkpoints have
passed
    without a new best
    exceeded ← patience > 0 and checkpoints_since_best >= patience
    # patience counts CHECKPOINTS (the plugin's own observable
cadence
    # from the log), not epochs – each script's own internal
epoch/val
    # cadence differs and doesn't matter here

```

```
write_best_checkpoint_pointer(models_dir, model_name, best_epoch):
    writes "<model_name>_best_recommended.txt" naming the
checkpoint file
    – a plain TEXT pointer, not a copy (checkpoints are 100s of
MB) and
    not an OS symlink (needs elevated privileges/Developer Mode on
Windows) – resolves identically cross-platform with zero
special
    permissions. Only needed for Cellpose-SAM, which has no built-
in
    best-tracking of its own (MONAI's train.py already auto-saves
its
    own best_model_fullstack.pth).
```

`_ai_tools.py` — argv builders

Pure argument-list construction (`build_prepare_data_argv /`
`build_monai_train_argv / build_cellpose_train_argv`) for the three
training scripts, which ship *inside the installed package*
(`training_scripts/`, declared as package data) rather than as sibling
folders in this project's private monorepo — a real historical bug: a
monorepo-sibling path only ever existed on the original dev machine,
so anyone else installing via `pip install git+https://...` got “script
not found” at launch. Package placement means they travel into site-
packages on both editable and normal installs. Always overridable via a
Browse button.

`_widget.py` — Tab 5 (AI Tools) handlers

GT Annotation section: buttons wrapping `interpolate_shapes()`/
`generate_masks()` directly — synchronous (both are fast: pure numpy/
shape rasterization, no model inference), errors shown in the status
label rather than crashing.

Extract Training Crops (two tools, side by side): the older bbox-based
single/double/triple/quadruple extractor (kept for anyone still using
that approach) and `generate_xyz_patches()` (the current default), both
background-thread wrapped with a cancel button (`cancel_event`) since
a full-fish run can take a while. “Clean crop truncation” is on by default
as part of the XYZY extractor's own run (not a separate button to
remember), per explicit instruction that Cellpose-SAM should train
only on near-complete cells ($\geq 90\%$) going forward.

Calibrate branch_radius (from GT): background-thread wrapper around `recommend_branch_radius()`, with the same one-shot “This is verified ground truth” gating pattern (§2/§5/§8) before it’s allowed to update the shared `branch_radius` config value.

Train MONAI U-Net / Train Cellpose-SAM buttons: build the argv via `_ai_tools.py`, call `launch_detached()`; poll `is_running/tail_log/patience_exceeded` on a `QTimer` while `napari` stays open, and **reconnect automatically on napari reopen** by reading the persisted PID + log path from config — a dead PID found on reopen finalizes immediately (best checkpoint reported, pointer written, stale PID cleared), so a job that finished while `napari` was closed is reported the moment `napari` reopens. “Stop Training” calls `kill_process_tree`. An optional Email-notification panel (SMTP fields, password via `_secrets.py`) threads a `notify dict` through to `launch_detached`.

8. Stage 6 — Sweeps & Utilities

UI: Tab 6 — Sweeps & Utilities. **Modules:** `_sieve.py` (generic coarse-to-fine engine, already covered in §2), `_brain_sweep.py`, `_pixel_sweep.py`, `_krendl_sweep.py`, `_epoch_sweep.py` (the four GT-verified parameter sweeps), `_gt_score.py` (the whole-fish scorer they build on), `_gt_package.py` (packaging for external manual correction), `_gt_toolkit.py` (single-folder auto-discovery tying them together). **Goal:** check a pipeline parameter against real ground truth and, once confirmed, recalibrate the “recommended” defaults quoted throughout §3–§5 — this is how those defaults were actually found, not guessed.

`_gt_score.py` — the whole-fish scoring methodology

`score_against_gt(pred_labels, gt_labels, iou_threshold=0.5)` is the project’s single, validated instance-matching methodology (used, in different guises, by every sweep in this section):

```
pred_info, gt_info ← per-object volume/centroid/bbox (_get_info
from _cellpose_seg.py)
build a (n_gt x n_pred) IoU cost matrix, but only for pairs whose
BBOXES
    actually intersect (IoU is necessarily 0 otherwise – keeps
this fast
    even for dozens of objects, without ever missing a real match)
row_ind, col_ind ← linear_sum_assignment(-iou_matrix)  #
Hungarian algorithm,
```

the GLOBALLY optimal one-to-one matching, not
a greedy

nearest-first assignment
for each assigned (gt, pred) pair:
 if IoU >= iou_threshold: count as TP, record IoU/Dice/
 size_delta
every unmatched GT object → FN; every unmatched pred object → FP
score ← TP - 0.5 × (FP + FN) # this project's
own established metric
mean_iou / mean_dice ← averaged over TP matches only

Verified against this project's own historical result tables (e.g. TP=34/
FP=8/FN=5 → Score=+27.5).

_brain_sweep.py — MONAI Threshold x Erosion (Tab 1's own mask)

The cheapest of the four sweeps: predict_probability() (the expensive network pass) runs **once**; every candidate (threshold, erosion) combination is then a cheap re-threshold via postprocess_probability() + a binary_erosion, scored as a whole-volume Dice/IoU/precision/recall against a hand-corrected GT brain mask (from GT Annotation, §7) — no per-cell matching needed, since brain segmentation is one global object.

```
pred_prob ← predict_probability(volume, model_path, device) #  
once, or  
    reused via precomputed= across sieve stages  
for each threshold:  
    raw_mask ← postprocess_probability(pred_prob, threshold)  
    for each erosion:  
        mask ← binary_erosion(raw_mask, erosion) if erosion > 0  
    else raw_mask  
        score whole-volume Dice/IoU/precision/recall vs GT  
best_point ← the (threshold, erosion) with highest Dice
```

_pixel_sweep.py — the Pixel Classifier route's two sweeps + GT-measured floors

Deliberately does **not** re-run MONAI inference per grid point (neither Signal Erosion nor BG Threshold change what MONAI predicted) — takes a pre-computed, un-eroded `brain_mask.tif` as input, making the whole sweep union-find-bound (minutes, no GPU needed) rather than GPU-inference-bound.

**run_pixel_sweep(image_path, brain_mask_path, gt_labels_path,
bg_thresholds, erosions, ...):**

```
cells ← find_complex_cells(gt_labels, n_cells)      # most skeleton
branches
    first, NOT by volume – a large cell is often simple/
amoeboid,
    while a genuinely branchy cell stresses the pipeline more
for each cell: crop image+brain_mask+GT to its own padded bbox
once
bg_max ← the GLOBAL background estimate (histogram over the whole
    volume's in-brain pixels) – computed once, since it's a
    genuinely global quantity, independent of Signal Erosion
for each bg_threshold:
    thresh ← bg_max + data_range x (bg_threshold / 100)
    for each erosion:
        for each cell's own crop:
            signal_mask ← erode_signal_2d((image > thresh) &
brain_mask, erosion)
                                # erodes what SURVIVED the threshold,
not the
                                # probe mask itself – matches
remove_global()
    pred_labels ← create_labels(image x signal_mask,
sigma_xy, sigma_z,
                                min_volume,
min_hole_size)
    score ← best-IoU match of any predicted object against
this cell's GT
    average IoU/Dice across all N cells at this grid point
best_point ← highest average IoU
```

run_sigma_sweep(...) — the Smooth σ_{XY}/σ_Z counterpart, roles reversed: BG Threshold/Signal Erosion held fixed, σ_{XY}/σ_Z swept. Cheaper per point than **run_pixel_sweep**, since σ only affects **create_labels()**, not the background-threshold crop-building step (computed once per cell, reused across every σ combination).

min_volume_from_gt(gt_labels): the smallest true voxel volume among GT-labeled cells, measured per fish rather than one constant shared across every fish regardless of how its real microglia actually run.

min_hole_size_from_gt(gt_labels, min_hole_size_to_trust=5): the smallest *real* internal hole any hand-corrected GT cell actually has (a human deliberately left it unlabeled) — holes below a trust threshold are discarded first, since real GT showed a sharp bimodal split (1–2 voxel annotation slips vs. a cluster of 400+ voxel real structural gaps, nothing in between); trusting every reported hole collapsed the recommendation to 0. **min_intercell_gap_um(gt_labels, scale_zyx, pad_um=15)**: smallest real physical surface-to-surface gap between any two distinct GT cells — a safety ceiling for Krendl safe-merge's **max_gap** (never bridge a gap that could be two real cells), computed per-cell-bbox rather than one full-volume distance transform (which timed out past 100s on a real fish).

_epoch_sweep.py — GT-verified Cellpose-SAM checkpoint sweep

Automates the bbox-restricted GT-IoU methodology developed by hand across this project's training rounds to confirm a checkpoint pick (from **test_loss**, a proxy metric) against real GT.

```
find_complex_cells(gt_labels, n_cells):
    rank every GT cell by (n_branches, -sphericity) – same
    skeleton/
    surface-area code as Statistics (§6), so this ranking matches
    what's
    reported elsewhere in the plugin
bbox_crop(image, gt_labels, label_id, pad_z=15, pad_xy=40): crop
to one
    cell's own padded bbox – do_3D cost scales with crop size, so
    this
    turns a full-fish sweep (~hours per checkpoint) into a small
    handful
    of seconds-scale crop inferences
```

```
run_epoch_sweep(image_path, gt_labels_path, models_dir,
```

```

model_name, epochs, ...):
    crop each of the N complex cells once
    for each candidate epoch:
        load that checkpoint ONCE (not once per cell)
        for each cell's crop: raw do_3D inference (no GMM/Krendl –
just
        best-IoU-match the single predicted object against
this cell's GT)
        average IoU/Dice per epoch across all N cells
best_epoch ← highest average IoU

pick_sweep_epochs(recommended_epoch, save_every, available_epochs,
n_below=2, n_above=2): centers the candidate epoch list on the
test_loss-recommended epoch, stepping by the training run's own
checkpoint interval, clipped to whatever actually exists on disk.

```

`_krendl_sweep.py` — Cellprob x Large-contact + real-data merge-parameter calibration

The most involved sweep module — three related but distinct calibration tools.

`run_krendl_sweep(volume, gt_labels, model_path, cellprobs, large_contacts, ...)` — the original instance-matched sweep, scored via `score_against_gt()` on the *fully corrected* result:

```

model, dP, cellprob_map, shape ← predict_flows(...) # ONE
network pass –
    cellprob_threshold only feeds the cheap masks_from_flows()
step
    (confirmed directly from cellpose/models.py's own eval()/
_compute_masks()
    split), so re-running the expensive GPU pass per cellprob
value would
    be needless; a multi-stage sieve narrowing can reuse this
same pass
    across every stage via precomputed=
for each cellprob:
    masks ← masks_from_flows(..., cellprob)
    masks ← gmm_cleanup(masks); masks ← krendl_safe_merge(masks,
max_gap,
                                min_contact, gt_min) # once per
cellprob
    for each large_contact: # cheap
on top of

```

```

        the same intermediate result
        merged ← large_contact_merge(masks, large_contact)
        merged ← final_min_size_cleanup(merged, gt_min,
final_min_fraction)
        labels ← relabel_sequential(merged)
        results[(cellprob, large_contact)] ←
score_against_gt(labels, gt_labels)
best_point ← highest Score, ties broken by mean_iou then
mean_dice, EXCLUDING
        any point with tp==0 (a point that detects nothing
can still
        "tie" a genuinely-detecting-but-noisy point on Score
– never a
        usable winner) unless literally every point detected
nothing

```

gt_min (Krendl safe-merge’s “already a whole cell” floor) and min_hole_size, if not given, are each measured from the real gt_labels passed in (min_volume_from_gt/min_hole_size_from_gt, imported from _pixel_sweep.py — the same measurement, since gt_min and the Pixel Classifier’s min_volume turned out to be literally the same quantity, unified after being tracked as two separate config histories by historical accident).

run_cellprob_voxel_sweep(volume, gt_labels, model_path, cellprobs, ...) — breaks a real circularity in the tool above: scoring cellprob against the *fully corrected* result entangles “best cellprob” with whatever max_gap/min_contact/large_contact happen to already be set to. This sweep scores cellprob on **pure voxel-level Dice/IoU** against binarized GT (instance identity ignored entirely) — mirrors _brain_sweep.run_brain_sweep()’s identical mask-level approach for MONAI Threshold. The documented correct calibration order: this sweep picks cellprob on signal quality alone → generate real cp_masks at that cellprob → the two functions below calibrate the merge parameters from that real prediction → optionally cross-check with run_krendl_sweep’s instance-matched Score using the now-non-circular calibrated parameters.

measure_merge_params_from_prediction(cp_masks, gt_labels, scale_zyx, min_overlap_vox=5, search_pad_um=3.0) — measures max_gap/min_contact from a real raw (pre-GMM, pre-Krendl) prediction rather than pure-GT geometry (which has a real ambiguity: is the smallest gap between two GT cells real biology, or just how an

annotator drew the shared boundary?). Sidesteps that by using GT only to *label* which raw fragments belong to which real cell, then measuring the network's own actual fragmentation:

```
assign every raw fragment to whichever GT cell it overlaps MOST
(drop as
    noise if the best overlap is < min_overlap_vox)
"should_merge" samples: for every GT cell with >= 2 assigned
fragments,
    each fragment's gap/contact to its NEAREST same-cell sibling
only –
    not all pairwise combinations (krendl_safe_merge() only ever
needs to
    bridge to the nearest neighbor, iteratively; all-pairs would
badly
    inflate the "should merge" gap distribution with meaningless
far-apart
    pairs within one large sprawling real cell)
"should_not_merge" samples: fragment pairs assigned to DIFFERENT
GT cells
    that end up close (bbox-proximity pre-filtered) – an
unambiguous
    safety-ceiling data point, since GT already confirms these are
genuinely separate cells
```

recommend_merge_params(merge_stats_list) — turns those pooled samples into one recommended (max_gap, min_contact) **pair**, jointly optimized against the REAL OR-combined decision rule (krendl_safe_merge merges if gap <= max_gap OR contact >= min_contact), not two independently-optimized single-criterion thresholds:

```
every OBSERVED (gap, contact) value becomes a candidate threshold
(exact,
    not a coarse grid)
for every candidate (max_gap, min_contact) pair:
    missed ← should-merge samples where BOTH gap > max_gap AND
contact < min_contact
        (computed for the whole grid in one matrix multiply,
        _count_gap_gt_and_contact_lt, not a triple-nested
loop)
    false ← should-not-merge samples where NEITHER criterion would
have
        excluded them (i.e. the pair WOULD wrongly merge
them)
```

```
total_error ← missed + false
best ← the (max_gap, min_contact) pair minimizing total_error
```

Optimizes for **fewest total manual corrections**, not zero risk of one error direction — an explicit, user-confirmed design choice: an over-merge is trivially fixable with Split Label, an under-merge with Join Labels (§5), so neither failure is silent or catastrophic; a defensively-clamped earlier version that minimized false-merges alone caught close to none of the real should-merge cases on real test data.

`_gt_package.py` — packaging for external manual correction

`build_gt_package(stem, out_dir, image_path, corrected_masks_path, raw_cellpose_masks_path=None, brain_mask_path=None, ...)` — automates a manual packaging step repeated by hand for every fish sent out for GT correction: copies the source image, the most-advanced correction stage available (as the correction starting point), the raw pre-merge masks (reference only), an optional brain mask (needed if the reviewer wants to use Protect Skin as Label themselves), and `GROUND_TRUTH_CREATION_GUIDE.md`, plus a lightweight per-cell CSV (`_cell_statistics_csv` — label/volume/centroid/bbox only, deliberately not the full ~51-column Tab 4 output), all zipped into `<stem>_GT_package.zip`.

`_gt_toolkit.py` — single-folder auto-discovery (“GT Toolkit Tuning Tool”)

Given one fish’s original source file (or its already-created output folder), locates every artifact this project’s established naming convention produces there and decides which of the above sweep/calibration steps can actually run without the user hunting down and browsing to each file by hand.

`SUFFIXES` ← the fixed map of {key: filename suffix} this plugin's own

```
    _output_dir() convention produces (_original.tif,
    _brain_mask.tif,
    _brain_only_ExtRm/NoBG/RndFill.tif, _cp.tif,
    _cp_krendl.tif,
    _cp_krendl_ac.tif, _cp_krendl_ac_snd.tif,
    _GROUND_TRUTH.tif,
    _statistics.csv) – Cellpose-SAM's own naming is
cumulative and
```

self-documenting: `_cp_krendl_ac_snd.tif`'s name itself says it went through Krendl, then auto-correct, then sanding

`discover_fish_files(source_path):`
`resolve_fish_folder(...)` ← accepts the folder itself, any file inside it, or the original raw source file sitting beside it
for each SUFFIXES key: an exact `<stem><suffix>` match, else the single result of a same-suffix glob if exactly one exists (tolerates the folder name and a file's own stem drifting apart slightly), else None

`best_corrected_masks(found):` the most-advanced stage present, priority
sanded > auto-corrected > Krendl-only (never the raw "cp" stage, which is always kept as a separate reference file)

`step_preconditions(found):` {step_key: (can_run: bool, reason: str)} for each of the 5 folded-in steps (monai/bg/sigma/cellprob/branch_radius)
– e.g. the MONAI sweep needs BOTH `_original.tif` AND `_brain_mask.tif` to exist, since their joint presence is what signals this fish went through GT Annotation (a raw MONAI-predicted `brain_mask.tif` alone isn't real ground truth to sweep against)

`_widget.py` — Tab 6 (Sweeps & Utilities) handlers

Every sweep tool shares: a **Sieve checkbox** (§2's `_add_sieve_controls()` — coarse full-range pass, then 1–2 narrowing stages around the winner) wired to `run_sieve()`; a one-shot **“This is verified ground truth”** checkbox (off by default, self-unticking after every run) gating whether the result is allowed to auto-apply to a live slider/config value or just reported; a **cancel button** wired to a `threading.Event` checked between grid points; and, once a sweep confirms a value, a call to

`_update_gt_history(config_key, fish_key, value, mode)` — the shared never-rising/never-falling cross-fish floor/ceiling mechanism every “recommended” default in this plugin is tracked through (§9).

GT Toolkit Tuning Tool: a single file/folder picker driving `discover_fish_files()` + `step_preconditions()`, pre-filling every other sweep tool’s own file fields with whatever it found and greying out (with a reason shown) any step this fish’s files can’t support.

Build GT-Correction Package: background-thread wrapper around `build_gt_package()`, auto-selecting `best_corrected_masks()` and the matching `_cp.tif/_brain_mask.tif` when a fish folder is picked via the GT Toolkit discovery above.

9. `_widget.py` — application shell

At 10,810 lines and 145 functions, `_widget.py` is both the biggest file in the plugin and the least algorithmically interesting one — the vast majority of it is straightforward Qt widget construction (`_build_ui()`, ~4,000 lines) and thin handler functions that read a few widgets, call into one of the modules covered in §3–§8, and write the result back to a napari layer via the threading pattern in §1. Rather than re-describe every handler a second time, this section covers the parts of the file that are genuinely its own infrastructure — construction, config persistence, cross-fish history, layer resolution, the hang watchdog, and email notification — and then indexes every handler by which pipeline-stage section already documents it.

Class structure

<code>_FastDirectLabelColormap(DirectLabelColormap)</code>	– module-level,
see §5	
<code>label")</code>	<code>("Hide skin</code>
<code>_load_config() / _save_config(data)</code>	– module-level,
plain	
<code>write of</code>	<code>JSON read/</code>
<code>napari-zf-</code>	<code>~/ .config/</code>
<code>config.json</code>	<code>microglia-ai/</code>
<code>_add_reliable_spinbox / _add_gt_checkbox /</code>	

```

_add_sieve_controls / _is_valid_cellpose_
checkpoint / _add_recommended_label /
_set_layout_widgets_visible / _make_collapsible /
_wrap_scroll / _sep / _make_notify_checkbox /
_extract_region_lines_um          – module-level
shared

                                UI helpers,
see §2/§6

```

```

class ZFMicrogliaAIWidget(QWidget):
    __init__(napari_viewer)
    _build_ui()          – constructs every tab's every widget
    _connect_signals()   – wires every button/combo/layer-event
to its handler
    <145 methods total>   – handlers, one section per tab, plus
shell helpers

```

__init__ and startup sequence

```

cfg ← _load_config()
cfg ← _secrets.migrate_plaintext_secrets(cfg)  # one-time
upgrade, see §2
model_path ← saved config path, else the bundled default
checkpoint, else
                None – resolved via is_file(), not exists(): Path("")
silently
                normalizes to Path(".") under pathlib, and exists()
is True
                for a directory too, which would otherwise let an
unconfigured
                path pass as "a valid model is loaded"
cellpose_model_path ← saved config only – no bundled default
(project-specific,
                not shipped with the plugin)
self._state ← {model_path, cellpose_model_path, last_file_path,
metadata, config}
self._skin_hidden_state ← None  # see _on_toggle_skin_visibility,
§5

_build_ui()
_connect_signals()          # also fires the initial-visibility
calls and
                                # _resume_monai_job_if_active() /
_resume_

```

```

# cellpose_job_if_active() ($7's
"reconnect
# on reopen" mechanism)
_refresh_layer_info()
_refresh_stats_layers()
_start_recovery_timer()

```

`_build_ui()` — organization, not algorithm

One long method building six `QTabWidget` pages in order (Skin Remover / Create MG Labels / Edit MG Labels / Statistics / AI Tools / Sweeps & Utilities), each wrapped in `_wrap_scroll()` and built from `_make_collapsible()` group boxes. There's no real control flow here worth pseudocode — every group follows the same shape: a `QGroupBox`, a `QVBoxLayout`, a slider+`_add_reliable_spinbox()` pair per numeric parameter, a status label, a run button. The one thing worth naming explicitly is the **Common Settings** group (Tab 2): Min volume and Min hole size are deliberately *read-only labels*, not editable sliders — both are empirical GT-measured facts (\$8's `min_volume_from_gt/` `min_hole_size_from_gt`), not tunable knobs a user should be able to accidentally drift away from a measured value.

Config persistence

```

_load_config() / _save_config(data):  plain JSON read/write, empty
dict / no-op
                                     on any failure – never
raises, so a
                                     corrupted or missing
config file never
                                     blocks the plugin from
opening
_save_cfg(self, **kwargs):  self._state["config"].update(kwargs);
_save_config(...)
                                     – the one method every handler
actually calls;
                                     merges rather than replaces, so
one handler's
                                     save never clobbers another's
unrelated keys

```


`_update_merge_stats_history(fish_key, merge_stats)` is the one exception, storing a whole *sample-list dict* per fish (not a scalar) and re-running `recommend_merge_params()` (§8) over the freshly-pooled set from every fish calibrated so far, rather than averaging a per-fish scalar.

Layer resolution

Three different “which layer” resolvers, each deliberately scoped to a different part of the plugin (see §1’s “Layer-selection philosophy”): - `_active_layer()`: active selection if it’s an Image layer, else the topmost Image layer, else None — the plugin’s original, still-used-for-Tab-1 fallback. - `_active_labels_layer()`: same pattern for Labels layers — used by a few route-agnostic tools (Split/Join/Save) that predate the Edit MG Labels tab’s shared selector. - `_edit_labels_layer()`: reads the *explicit* combo the user set in Edit MG Labels’ own “Layers and label(s) being edited” section — every Stage-3 tool (§5) uses this one, never a fallback, so the layer being edited is never guessed.

`_get_layer_scale()`: metadata from the last Open-button load, else the active layer’s own napari scale, else (1,1,1) — the one function nearly every physical-unit calculation in the plugin ultimately calls.

`_output_dir()`: <original file's parent>/<original file's stem>, created on first use — the one folder convention every saved output file in the plugin follows (and that §8’s `_gt_toolkit.py` reverse-engineers to auto-discover a fish’s files). `_current_min_volume()` / `_lo_sweep_sample_params()`: thin config readers so every route/sweep launcher shares one source of truth instead of each hardcoding its own possibly-diverging copy.

The hang watchdog

```
_hang_watch_start(what, after_s=600):
    faulthandler.dump_traceback_later(after_s, repeat=True,
file=sys.stderr)
    # if the operation is STILL running after after_s seconds (and
every
    # after_s after that), every thread's real stack prints to the
    # terminal napari was launched from – turns "is it stuck?"
into
    # "here's exactly which line it's on", instead of only the
last
    # status message (which can stay unchanged through several
later steps)
```

```
_hang_watch_stop():  
    faulthandler.cancel_dump_traceback_later()
```

Wrapped around Protect Skin as Label, Auto-correct, and the Cellpose-SAM-chained auto-correct stage — the slowest operations in the plugin, extended to 30 minutes (from an original 10) after real full-fish runs were found to legitimately take that long.

Email notification

Two independent paths sharing one credential source (`_get_notify_creds()`, reading Tab 6's shared SMTP fields and persisting the password via `_secrets.py`, never plaintext): - **In-process** (`_maybe_send_notify(checkbox, subject, body)`): called from a background *worker* thread right before it finishes (not the GUI/poll thread, so the SMTP round-trip never blocks the UI) — used by Tab 1 Run, Tab 2 Cellpose-SAM Segmentation, and Tab 6's Cellprob/Large-contact and Best-EPOCH sweeps, none of which survive napari closing anyway. - **Detached** (`_build_notify_cfg() → _training_jobs.launch_detached(..., notify=...)`, §7): used only by Tab 5's two training launchers, which must keep running (and therefore be able to notify) even after napari itself has closed — hence the standalone supervisor-script approach instead of this same in-process function. Both paths swallow their own errors into a `print()` rather than raising — a broken email config must never take down the operation whose completion it was reporting.

The 10-minute recovery timer

`_start_recovery_timer()` — the one `QTimer` in the plugin meant to live for the whole session (every other background-work `QTimer` calls `deleteLater()` once its job finishes; a real historical OOM leak came from `QTimers` that didn't, keeping their connected closures — often a full label-array copy — pinned in memory for the rest of the session). Every 10 minutes, saves every current `Labels` layer to `<output_dir>/<layer_name>_recovery.tif` in a background thread, overwriting the previous save — a crash loses at most ~10 minutes of manual label editing, not the whole session.

`_refresh_*` family — keeping the UI in sync with the viewer

Each wired to `viewer.layers.events.inserted/.removed` (and, for the active-selection-dependent ones, `.selection.events.changed`), so the UI never goes stale as layers are added/removed/renamed elsewhere: - `_refresh_layer_info()` / `_update_labels_section_visibility()`: shows/hides the Pixel Classifier vs. Cellpose-SAM Segmentation groups based on whether a matching `_NoBG/_ExtRm`-suffixed Image layer *exists anywhere in the viewer* — not just whether it’s the currently active layer (a real historical bug: the old active-layer-only check could hide a route’s tools the moment the user selected the Labels layer itself, exactly what “Use selected” requires). - `_refresh_stats_layers()`: repopulates Tab 4’s Image/Shapes combos (§6). - `_refresh_gt_layers()`: repopulates Tab 5’s GT Annotation Image combo (§7). - `_refresh_meta_lbl()`: updates the Z/Y/X μm + anisotropy display under Tab 1’s Open button.

Handler index — every `_widget.py` function, by home section

Every button/combo handler in the file is documented, with real behavior, in the pipeline-stage section for the tab it belongs to — listed here so nothing in the file is left unaccounted for:

Tab	Section	Representative handlers
1 — Skin Remover	§3	<code>_on_open</code> , <code>_on_load_labels</code> , <code>_on_run</code> , <code>_on_browse_model</code> , <code>_on_bg_mode_changed</code> <code>_on_create_labels</code> , <code>_on_run_cellpose_seg</code> , <code>_on_rerun_single_cell</code> , <code>_resolve_rerun_layers</code> , <code>_on_browse_cp_model</code> , <code>_run_auto_correction_stage</code> , <code>_run_sanding_stage</code> <code>_on_resort_labels</code> , <code>_on_remove_debris</code> , <code>_on_split_label</code> , <code>_on_join_labels</code> , <code>_on_correct_label</code> , <code>_on_correct_adjacent_labels</code> , <code>_on_copy_label_to_adjacent_slice</code> , <code>_on_protect_skin</code> ,
2 — Create MG Labels	§4	
3 — Edit MG Labels	§5	

Tab	Section	Representative handlers
		_on_remove_skin_label, _on_toggle_skin_visibility, _on_autocorrect_labels, _on_save_labels, _on_run_contrast_sweep, _on_toggle_drift, _on_use_selected_label_a/b/rerun, _on_correct_label_id_changed
4 — Statistics	§6	_on_generate_stats, _on_stats_backend_changed, _on_gtscore_run
		_on_gt_interpolate, _on_gt_generate_masks, _on_gt_image_changed, _on_prepare_monai_data, _on_mt_launch_training, _on_mt_stop_training, _resume_monai_job_if_active, _start_monai_polling, _on_ct_launch_training, _on_ct_stop_training, _resume_cellpose_job_if_active, _start_cellpose_polling, _on_ct_calib_run (branch radius), _on_es_run_sweep (epoch sweep), _on_xz_run (XZYZ patches), _write_cellpose_best_pointer
5 — AI Tools	§7	
		_on_gtk_scan/_on_gtk_run (GT Toolkit), _on_bs_run_sweep (brain/MONAI), _on_ps_run_sweep (BG Threshold/Erosion), _on_sg_run_sweep (Sigma), _on_kr_run_sweep (Cellprob/Large-contact), _on_gtp_run (GT package)
6 — Sweeps & Utilities	§8	
		__init__, _build_ui, _connect_signals, _status, _save_cfg, _update_gt_history, _update_merge_stats_history, _hang_watch_start/stop, _get_notify_creds, _build_notify_cfg, _maybe_send_notify,
Shell (this section)	§9	

Tab	Section	Representative handlers
		_on_send_test_email, _start_recovery_timer, _output_dir, _get_layer_scale, _active_layer, _active_labels_layer, _edit_labels_layer, _refresh_*, _current_min_volume, _lo_sweep_sample_params, _on_ai_tools_mode_changed, _update_labels_section_visibility

10. Module reference (quick index)

File	Lines	Purpose	Guide section
_widget.py	10,810	Application shell — all 6 tabs, every button handler, config, layer resolution	§3–§9
_labeling.py	2,883	Pixel Classifier connected-components engine + every label-editing primitive (resort, debris, split, join, correct, copy, skin protection)	§4, §5
_auto_correction.py	993	6-step unattended per-cell auto-correction pipeline, chained after Cellpose-SAM Segmentation	§5
_cellpose_seg.py	808	Cellpose-SAM do_3D inference + GMM cleanup + Krendl safe-merge + large-contact merge + final safety net	§4
_statistics.py	764		§6

File	Lines	Purpose	Guide section
		Per-label morphology/ spatial/intensity statistics + natural- language descriptions	
<code>_krendl_sweep.py</code>	719	Cellprob x Large- contact GT sweep + real-prediction merge- parameter calibration	§8
<code>_pixel_sweep.py</code>	609	BG Threshold x Signal Erosion / Smooth σ_{XY} x σ_Z GT sweeps + GT- measured floors (<code>min_volume_from_gt</code> , <code>min_hole_size_from_gt</code> , <code>min_intercell_gap_um</code>)	§8
<code>_io.py</code>	284	Metadata-aware .tif/.ims loading, voxel-scale resolution	§2
<code>_grow_correct.py</code>	493	Auto-grow / until- stable orchestration around Correct Label's 2D and 3D engines	§5
<code>_training_jobs.py</code>	378	Cross-platform detached-process training job management (launch, poll, kill, patience early-stop, email)	§7
<code>_gt_annotation.py</code>	322	Polygon-based hand- drawn GT annotation (brain/skin masks)	§7
<code>_epoch_sweep.py</code>	275	GT-verified Cellpose- SAM checkpoint (epoch) sweep, bbox- restricted	§8
<code>_contrast_sweep.py</code>	275	Self-referential Correct-Label contrast (1o) calibration against	§5

File	Lines	Purpose	Guide section
		Cellpose-SAM's own output	
<code>_background.py</code>	193	Population-mode background estimation + 3 background-removal modes	§3
<code>_xyz_patches.py</code>	187	Orientation-slice (XY/XZ/YZ) 2D training-crop generation, the current Cellpose-SAM training methodology	§7
<code>_gt_toolkit.py</code>	180	Single-folder auto-discovery of a fish's files + step preconditions ("GT Toolkit Tuning Tool")	§8
<code>_gt_package.py</code>	167	Zips a GT-correction package (source image, correction starting point, guide, stats CSV) for external manual review	§8
<code>_brain_sweep.py</code>	165	GT-verified MONAI Threshold x Erosion sweep for the brain mask	§8
<code>_secrets.py</code>	188	Two-tier (OS keyring → local Fernet-encrypted file) secret storage	§2
<code>_crop_truncation.py</code>	146	Zeros out incidentally-truncated neighbor labels in training crops	§7
<code>_branch_calibration.py</code>	134	Measures real branch radius from GT to recalibrate the branch-weighted training loss	§7
<code>_gt_score.py</code>	126		§8

File	Lines	Purpose	Guide section
<code>_live_progress.py</code>	116	Whole-fish Hungarian-matched TP/FP/FN/Score/IoU/Dice scoring — the shared methodology every sweep builds on Captures MONAI's/Cellpose's own otherwise-invisible progress output into the GUI log	§2
<code>_sieve.py</code>	105	Generic coarse-to-fine sweep-narrowing engine, shared by every continuous-axis GT sweep	§2
<code>_ai_tools.py</code>	89	argv builders + default script-path resolution for the 3 training scripts	§7
<code>_sanding.py</code>	134	Sigma-softening (contour polishing) pass, chained after any label-regenerating correction	§4
<code>_reader.py</code>	58	napari File-menu integration (.tif/.ims reader registration)	§2
<code>_gpu_check.py</code>	55	GPU/VRAM classification (informational banner, never a hard gate)	§2
<code>__main__.py</code>	53	CLI entry point / pre-loading a file at launch	—
<code>_inference.py</code>	169	MONAI 3D U-Net sliding-window brain inference	§3
<code>__init__.py</code>	22	Package version + widget export	—

(training_scripts/ — prepare_data.py, train.py, train_xyz.py — *ships inside the package as launchable subprocess scripts, not imported; see _ai_tools.py and §7.*)

End of Software Guide.